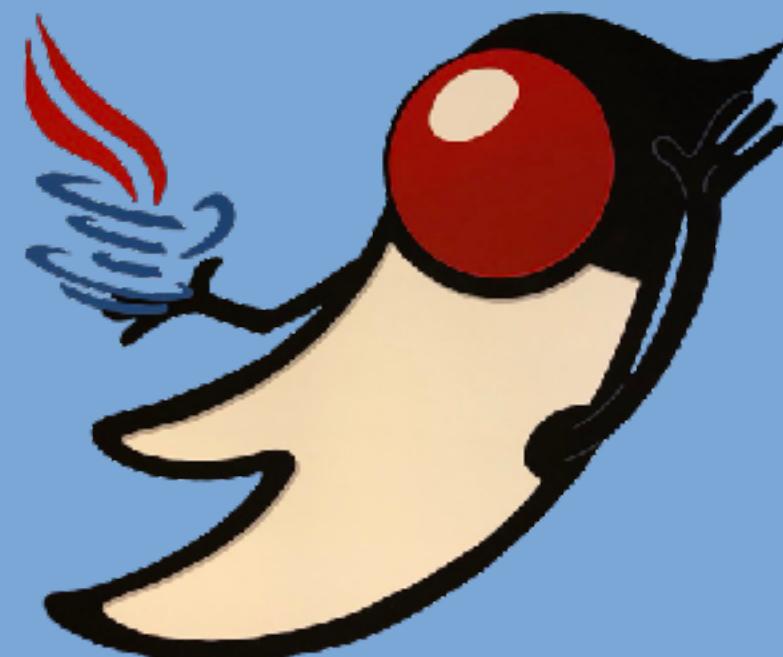


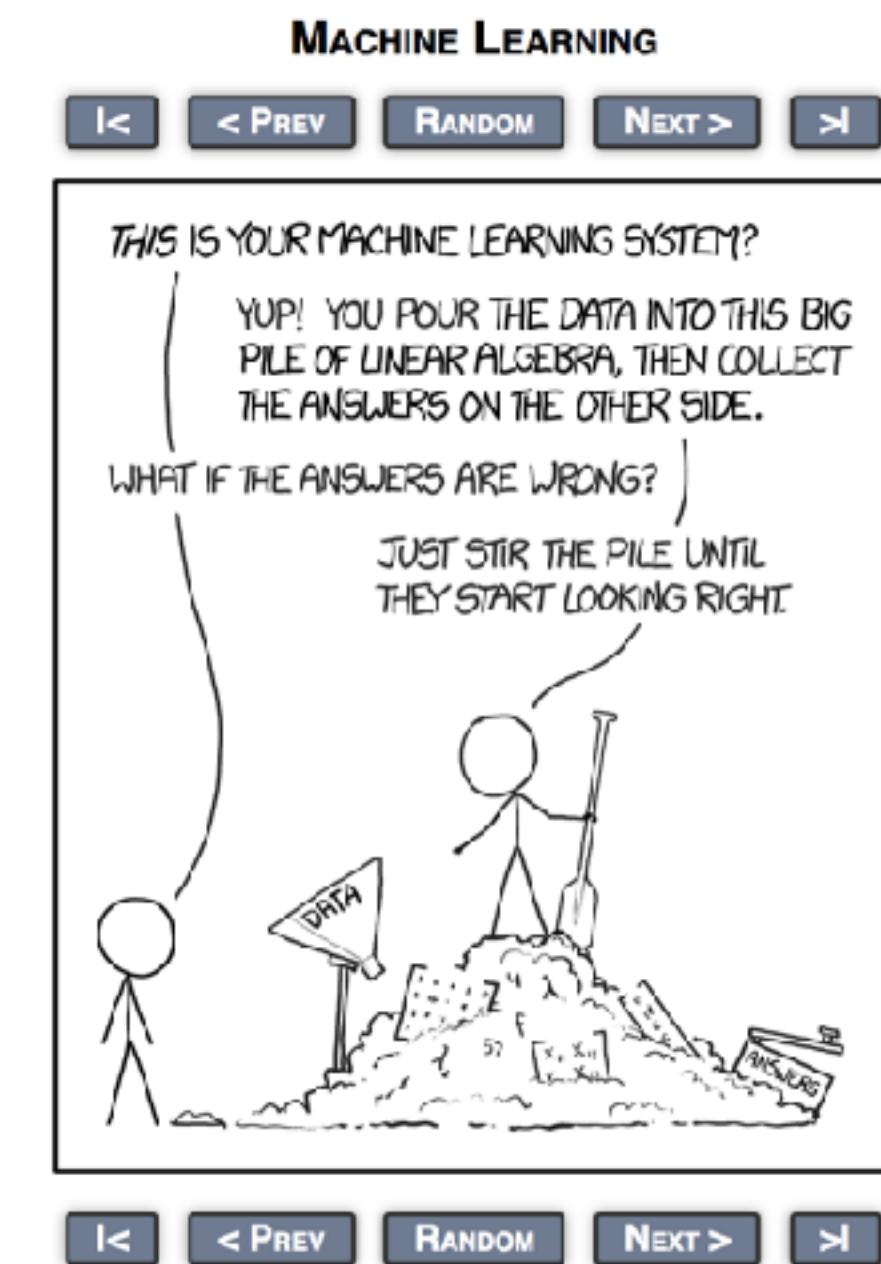
Bayesian Optimization of Gaussian Processes applied to Performance Tuning

Ramki Ramakrishna
@ysr1729
#TwitterVMTeam

QCon Sao Paulo, 2019



A JVM Engineer talks to a Data Scientist



PERMANENT LINK TO THIS COMIC: <https://xkcd.com/1838/>

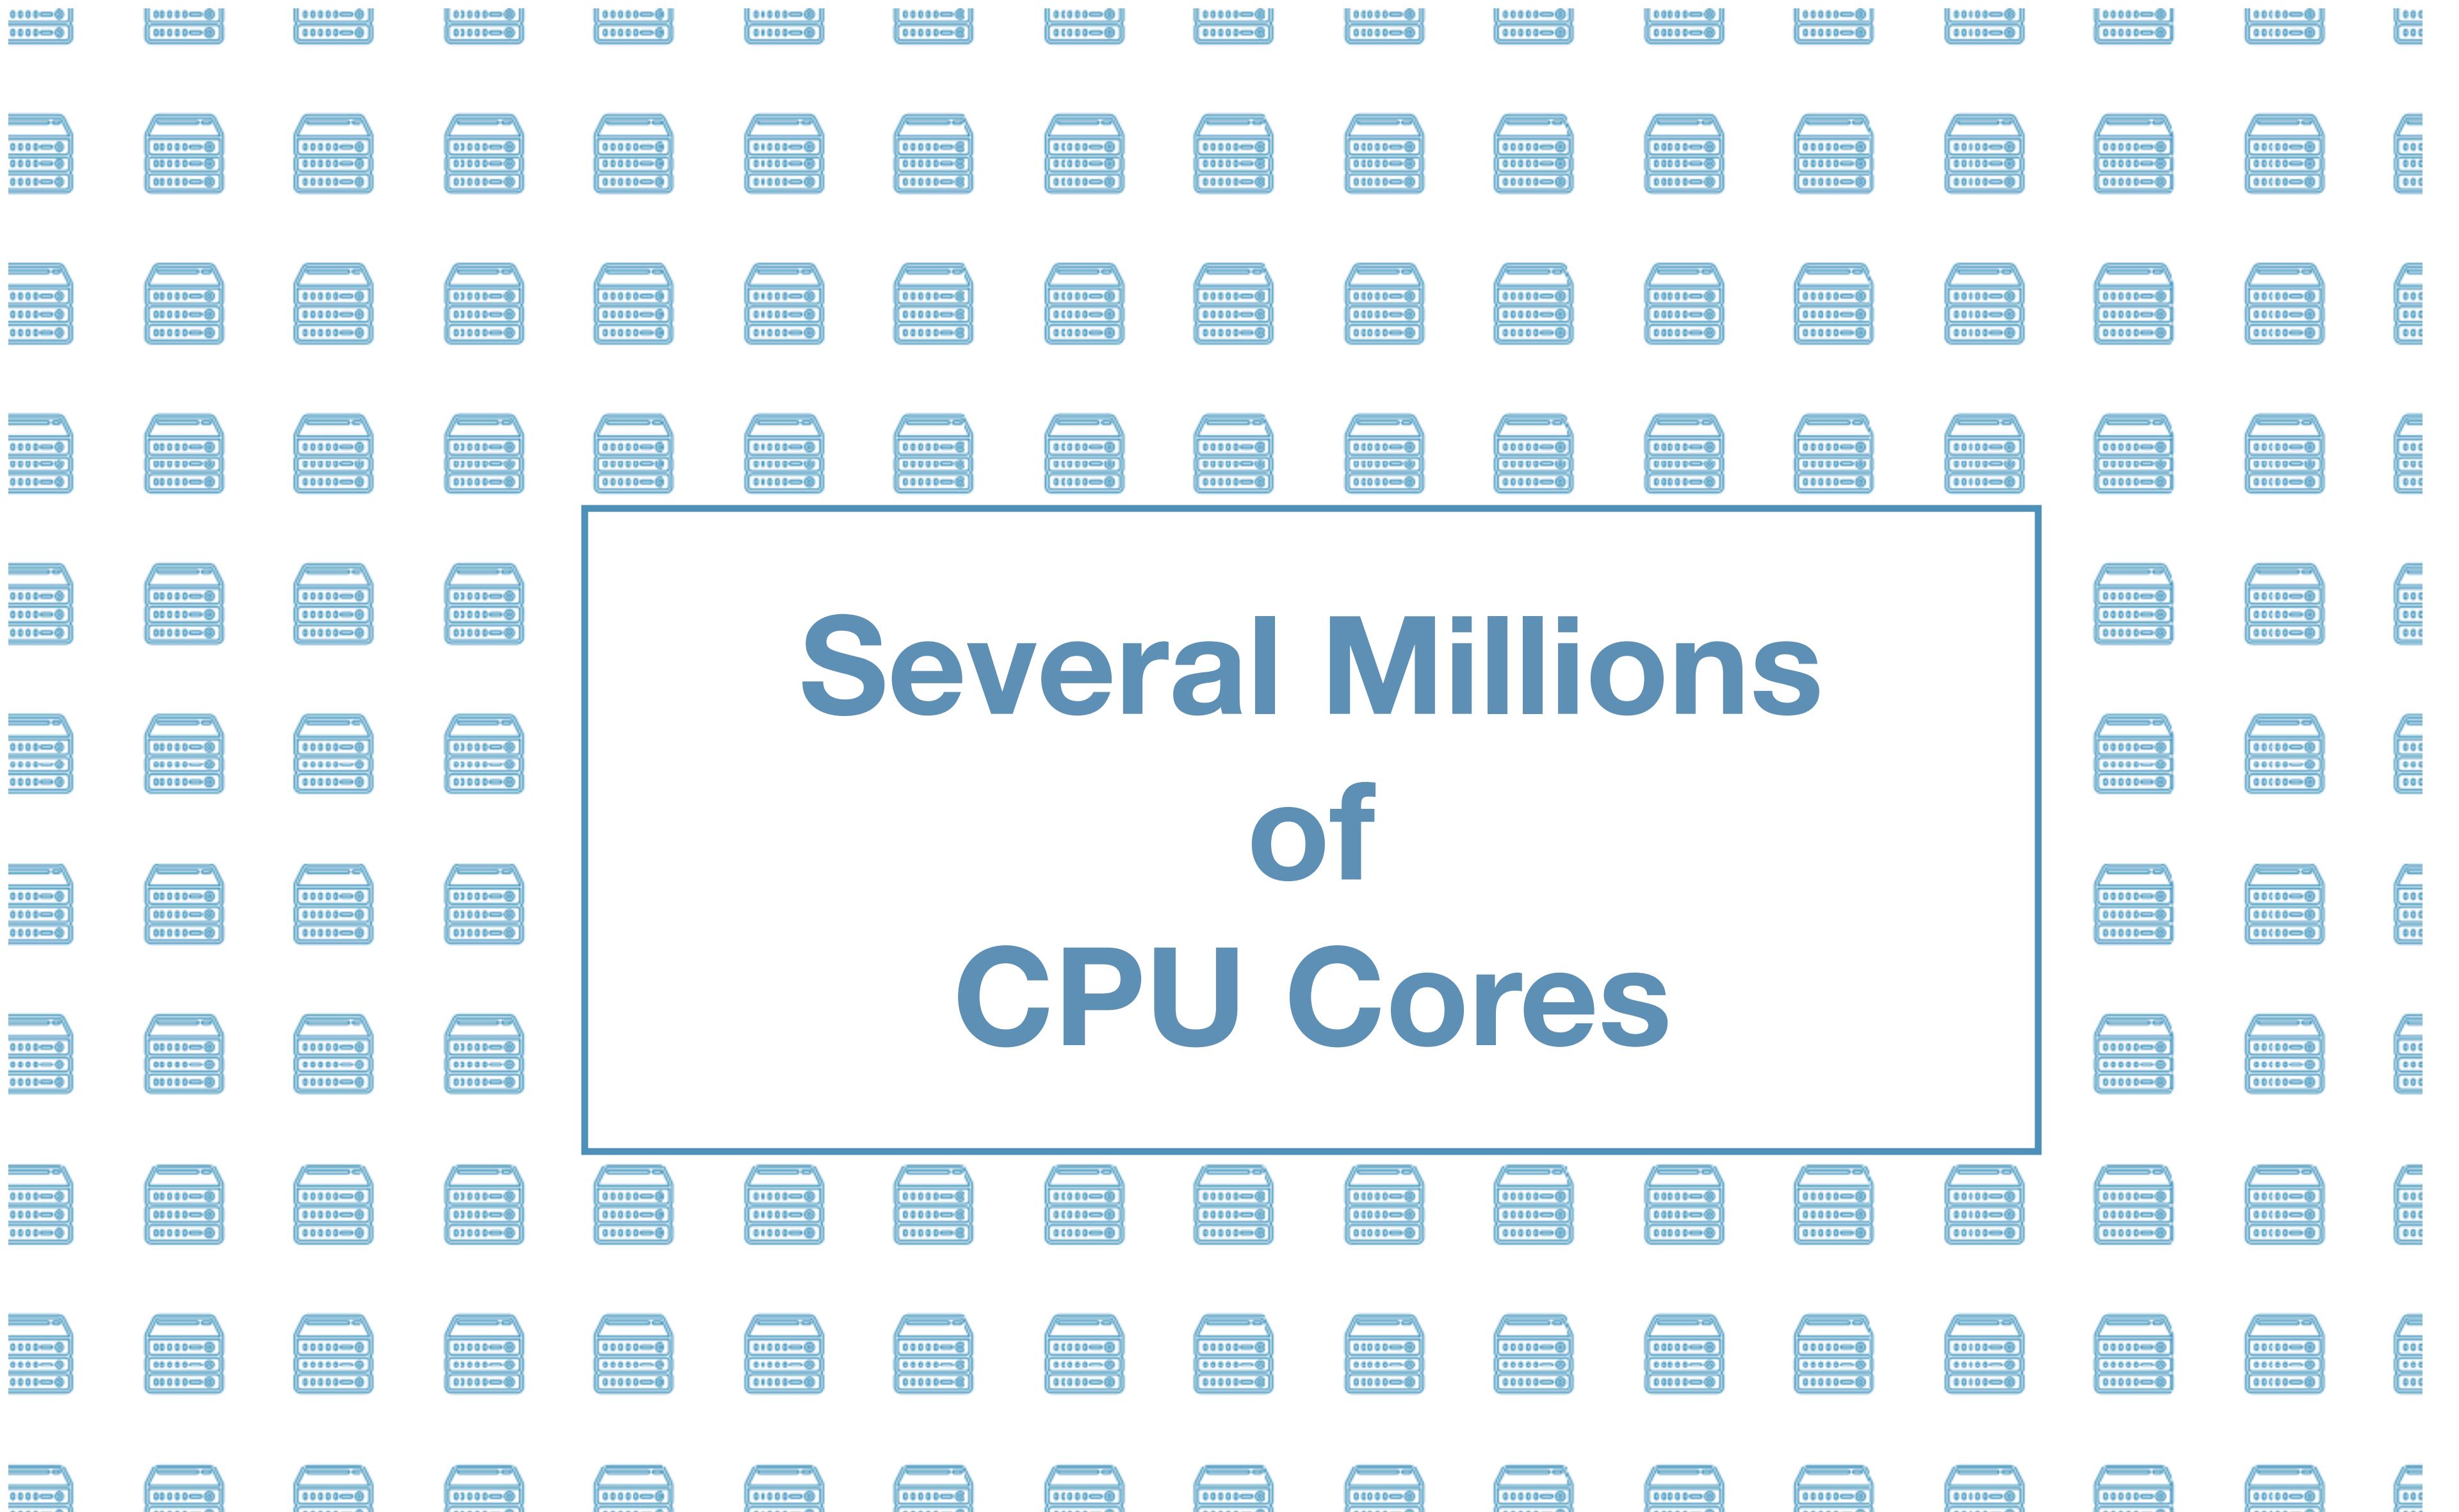




**Many Hundreds
of
Services**



**Several Tens of
Thousands of
Physical Servers**



**Several Millions
of
CPU Cores**



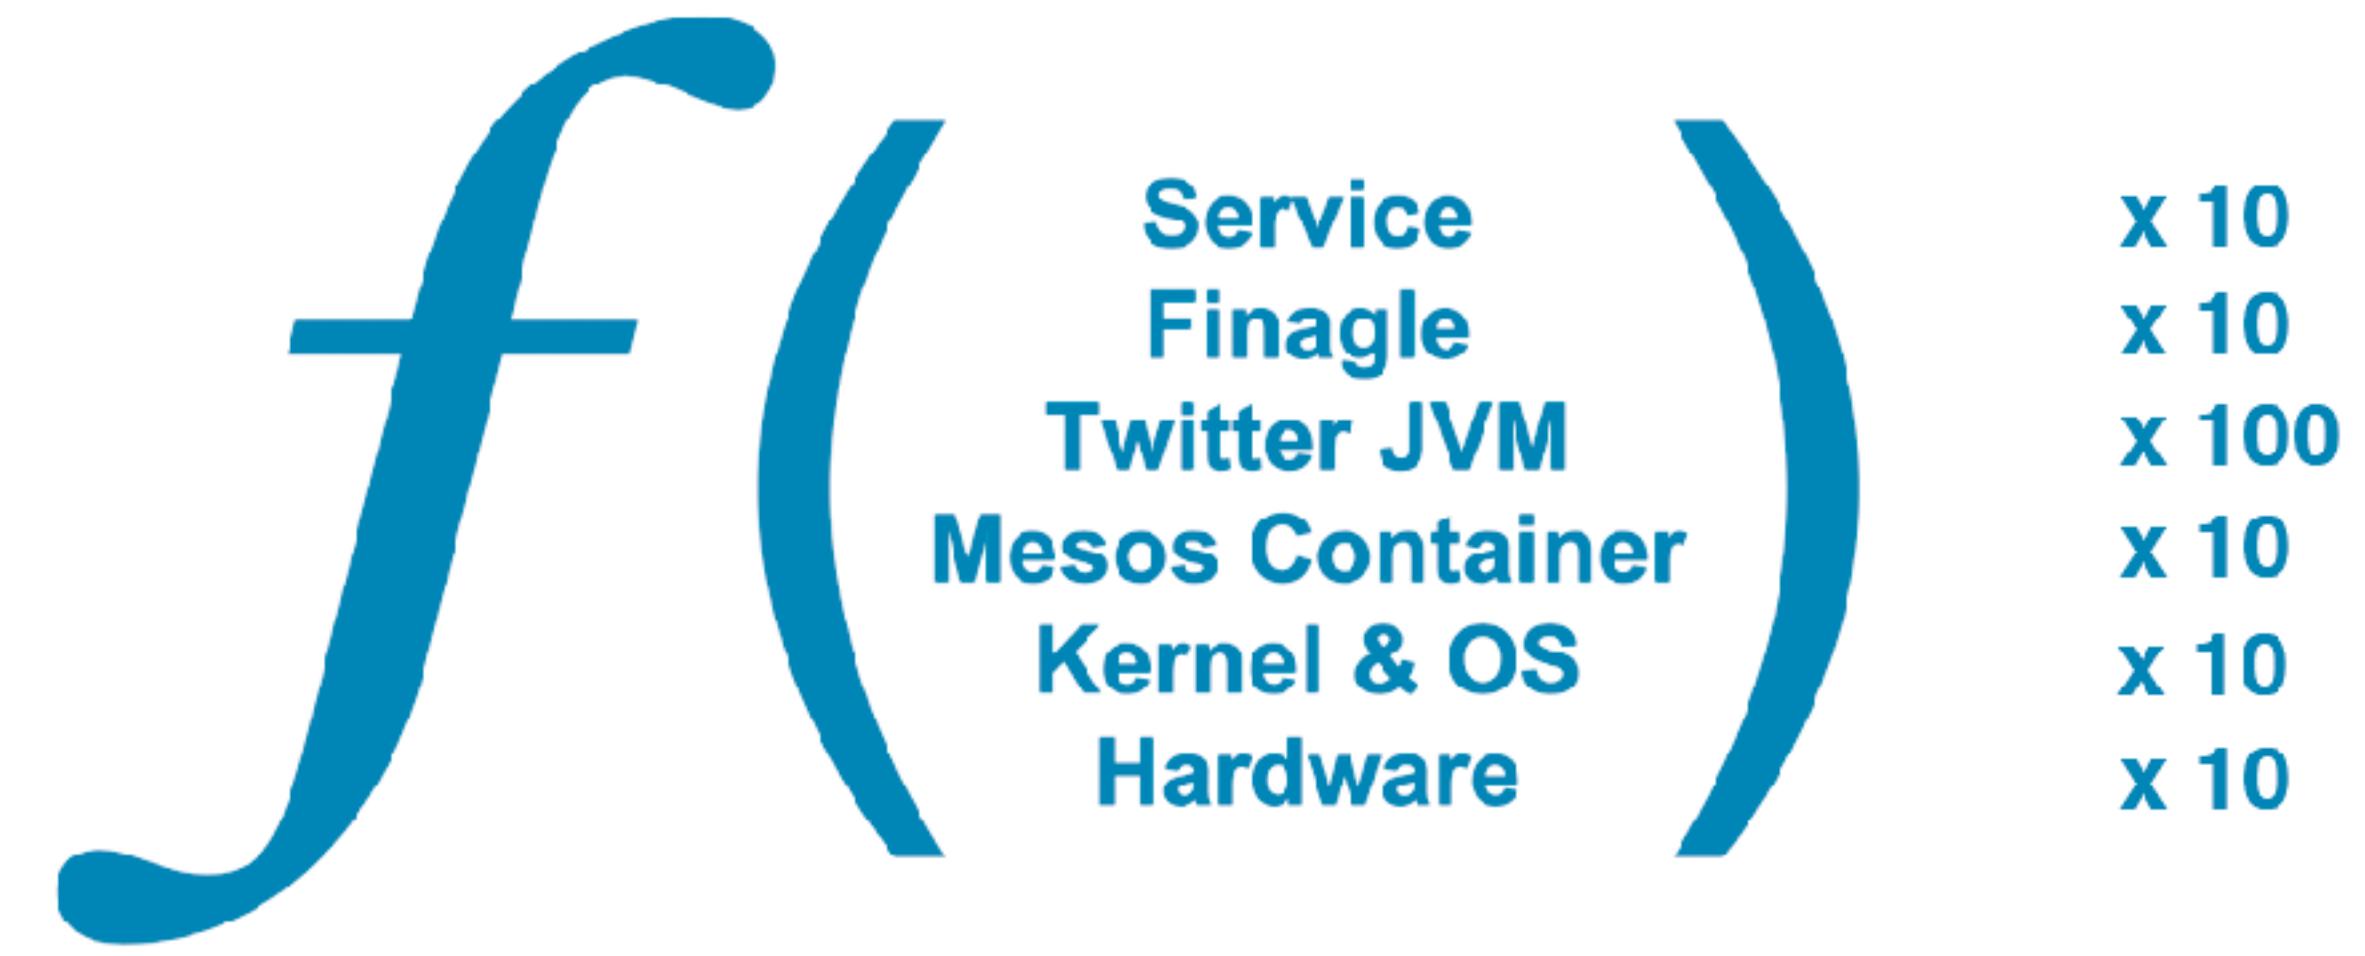
**Several Hundreds of
Thousands of Twitter
JVMs**



A Few Hundred Tunable JVM Parameters







x 10
x 10
x 100
x 10
x 10
x 10





Andrew Ng

@AndrewYNg

If you're trying to understand AI's near-term impact, don't think "sentience." Instead think "automation on steroids."

11:04 AM - 1 May 2017 from Stanford, CA

1,397 Retweets 2,085 Likes



38

1.4K

2.1K



Mining for Gold

- 1930's South Africa
- Prospecting for gold and other minerals
- Daniel Krige, 1951: “Kriging” in geostatistics
- Jonas Mockus 70's
- Jones et al. 80's
- Rasmussen & Williams 90's : Gaussian Processes



Applications

- design of expensive experiments
- optimal designs
- optimization of engineered materials
- hyperparameter tuning (architectural parameters) of neural networks

Engineering as Optimization

- linear or non-linear objective function
- finite convex or non-convex space, rectangular, linear (affine) or non-linear constraints
- black box objective function
- black box constraints
- noisy objective function
- noisy constraints



Black Box Modeling

- Model the unknown objective function
- Model the unknown constraints
- Model is a “surrogate”
- Evaluations are expensive



Models and Model Parameters

- Parametric models
- Non-parametric models



Probabilistic Models

- A measure of our uncertainty
- A measure of measurement/observation noise



Gaussian Process

$$GP(\mu, \kappa)$$

- $\mu(x)$: mean function
- $\kappa(x, x')$: covariance function

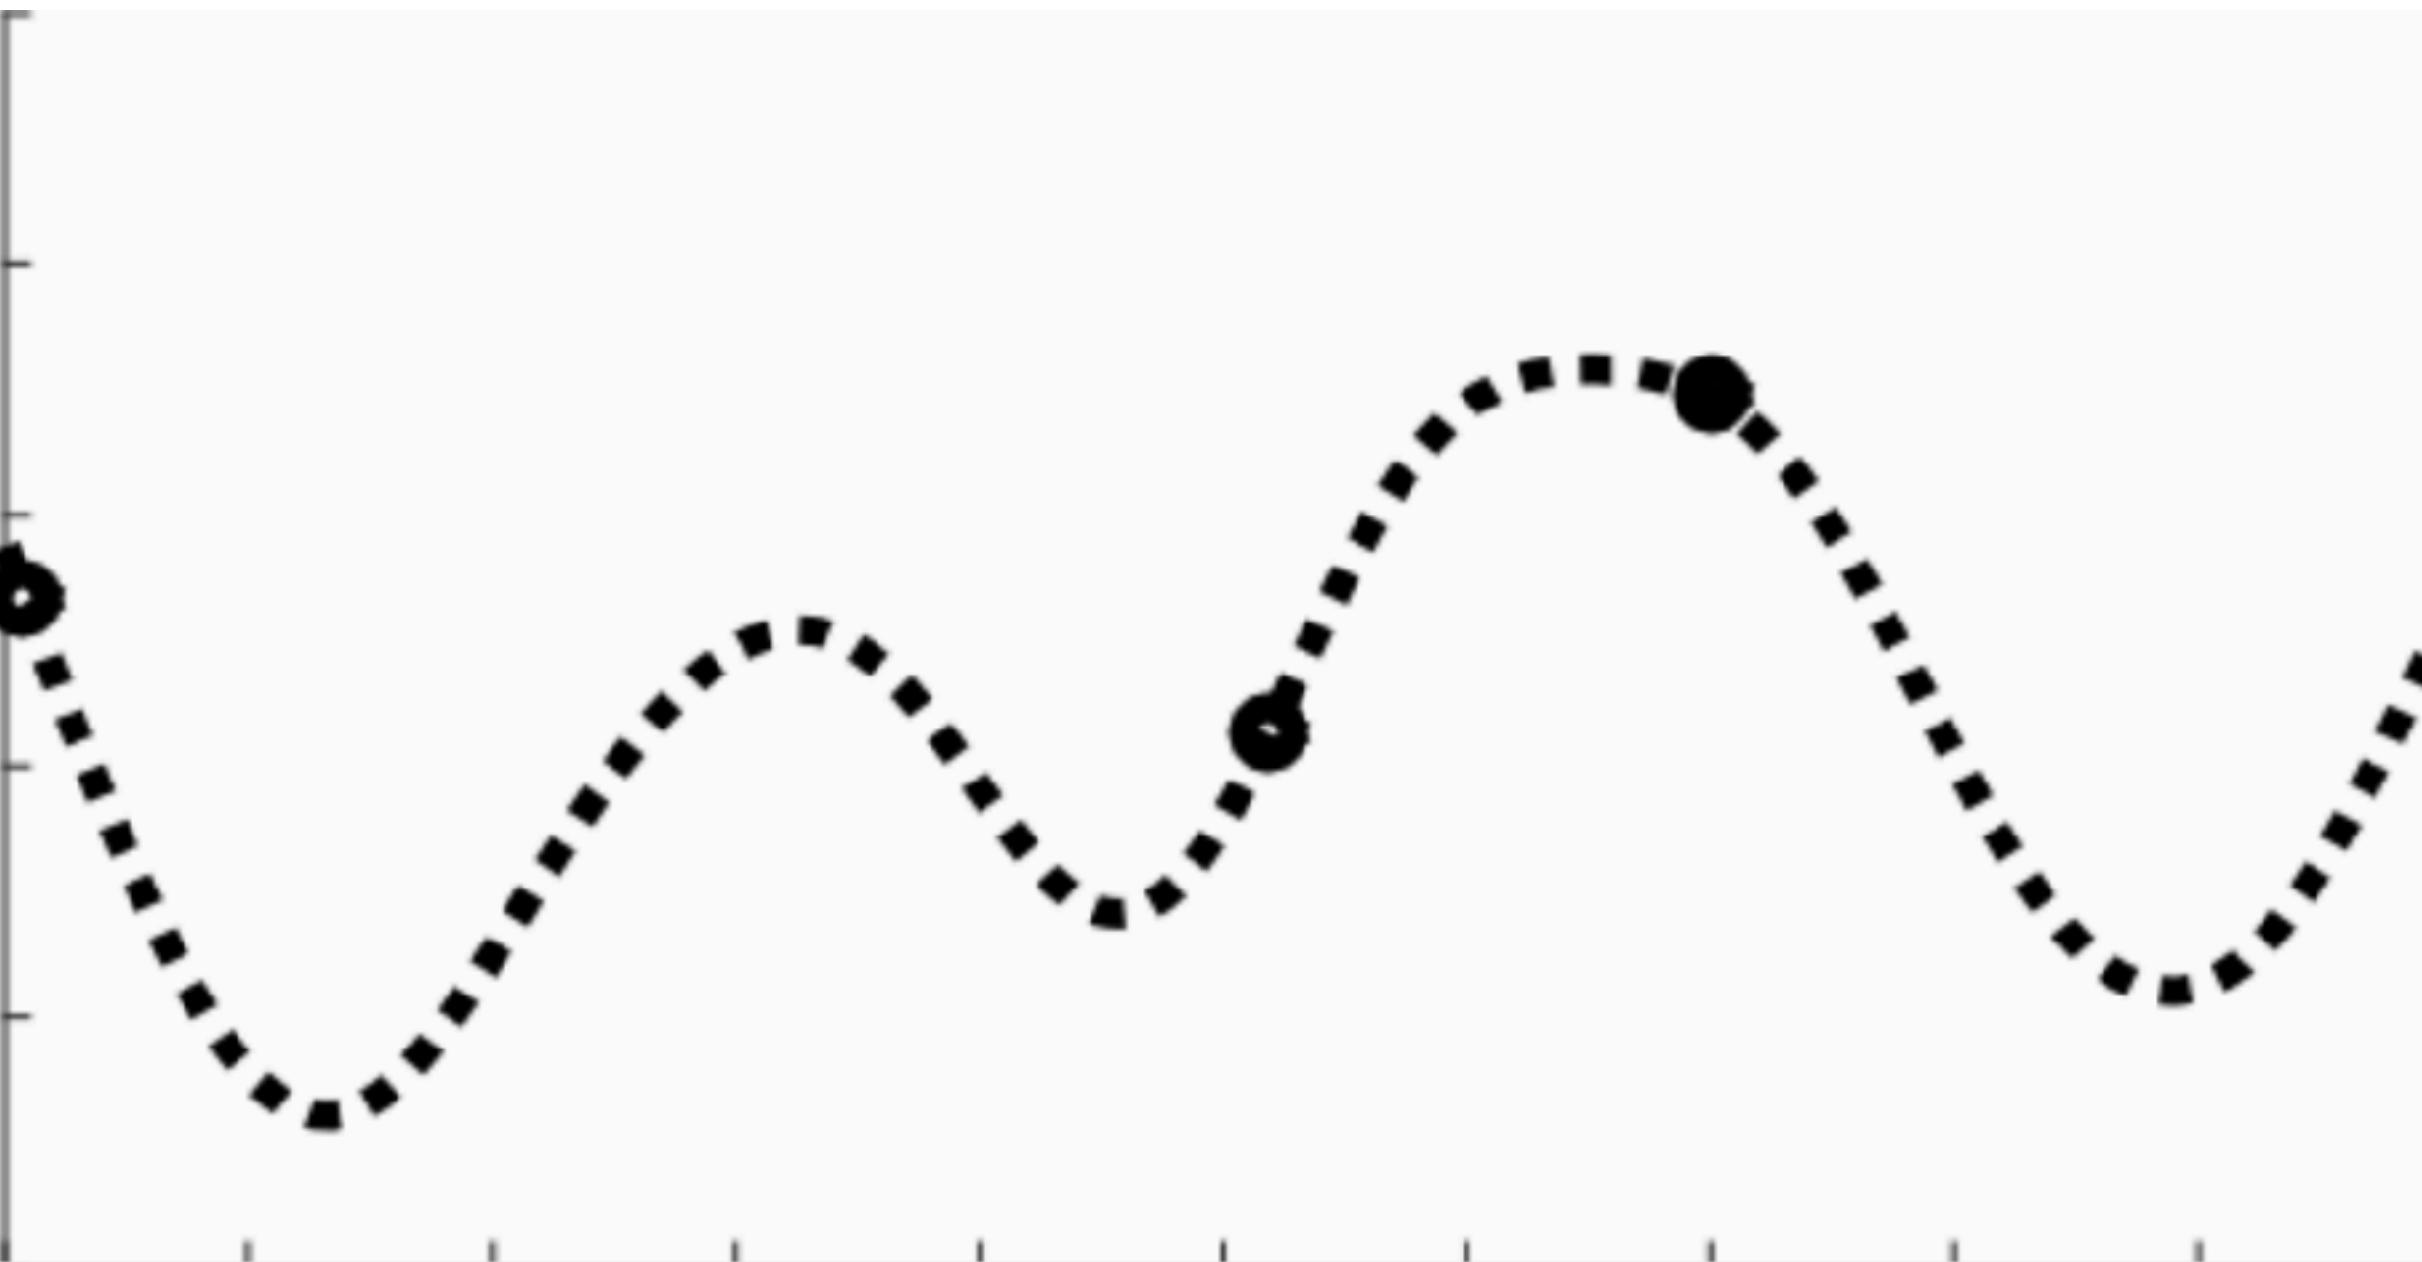


Gaussian Process

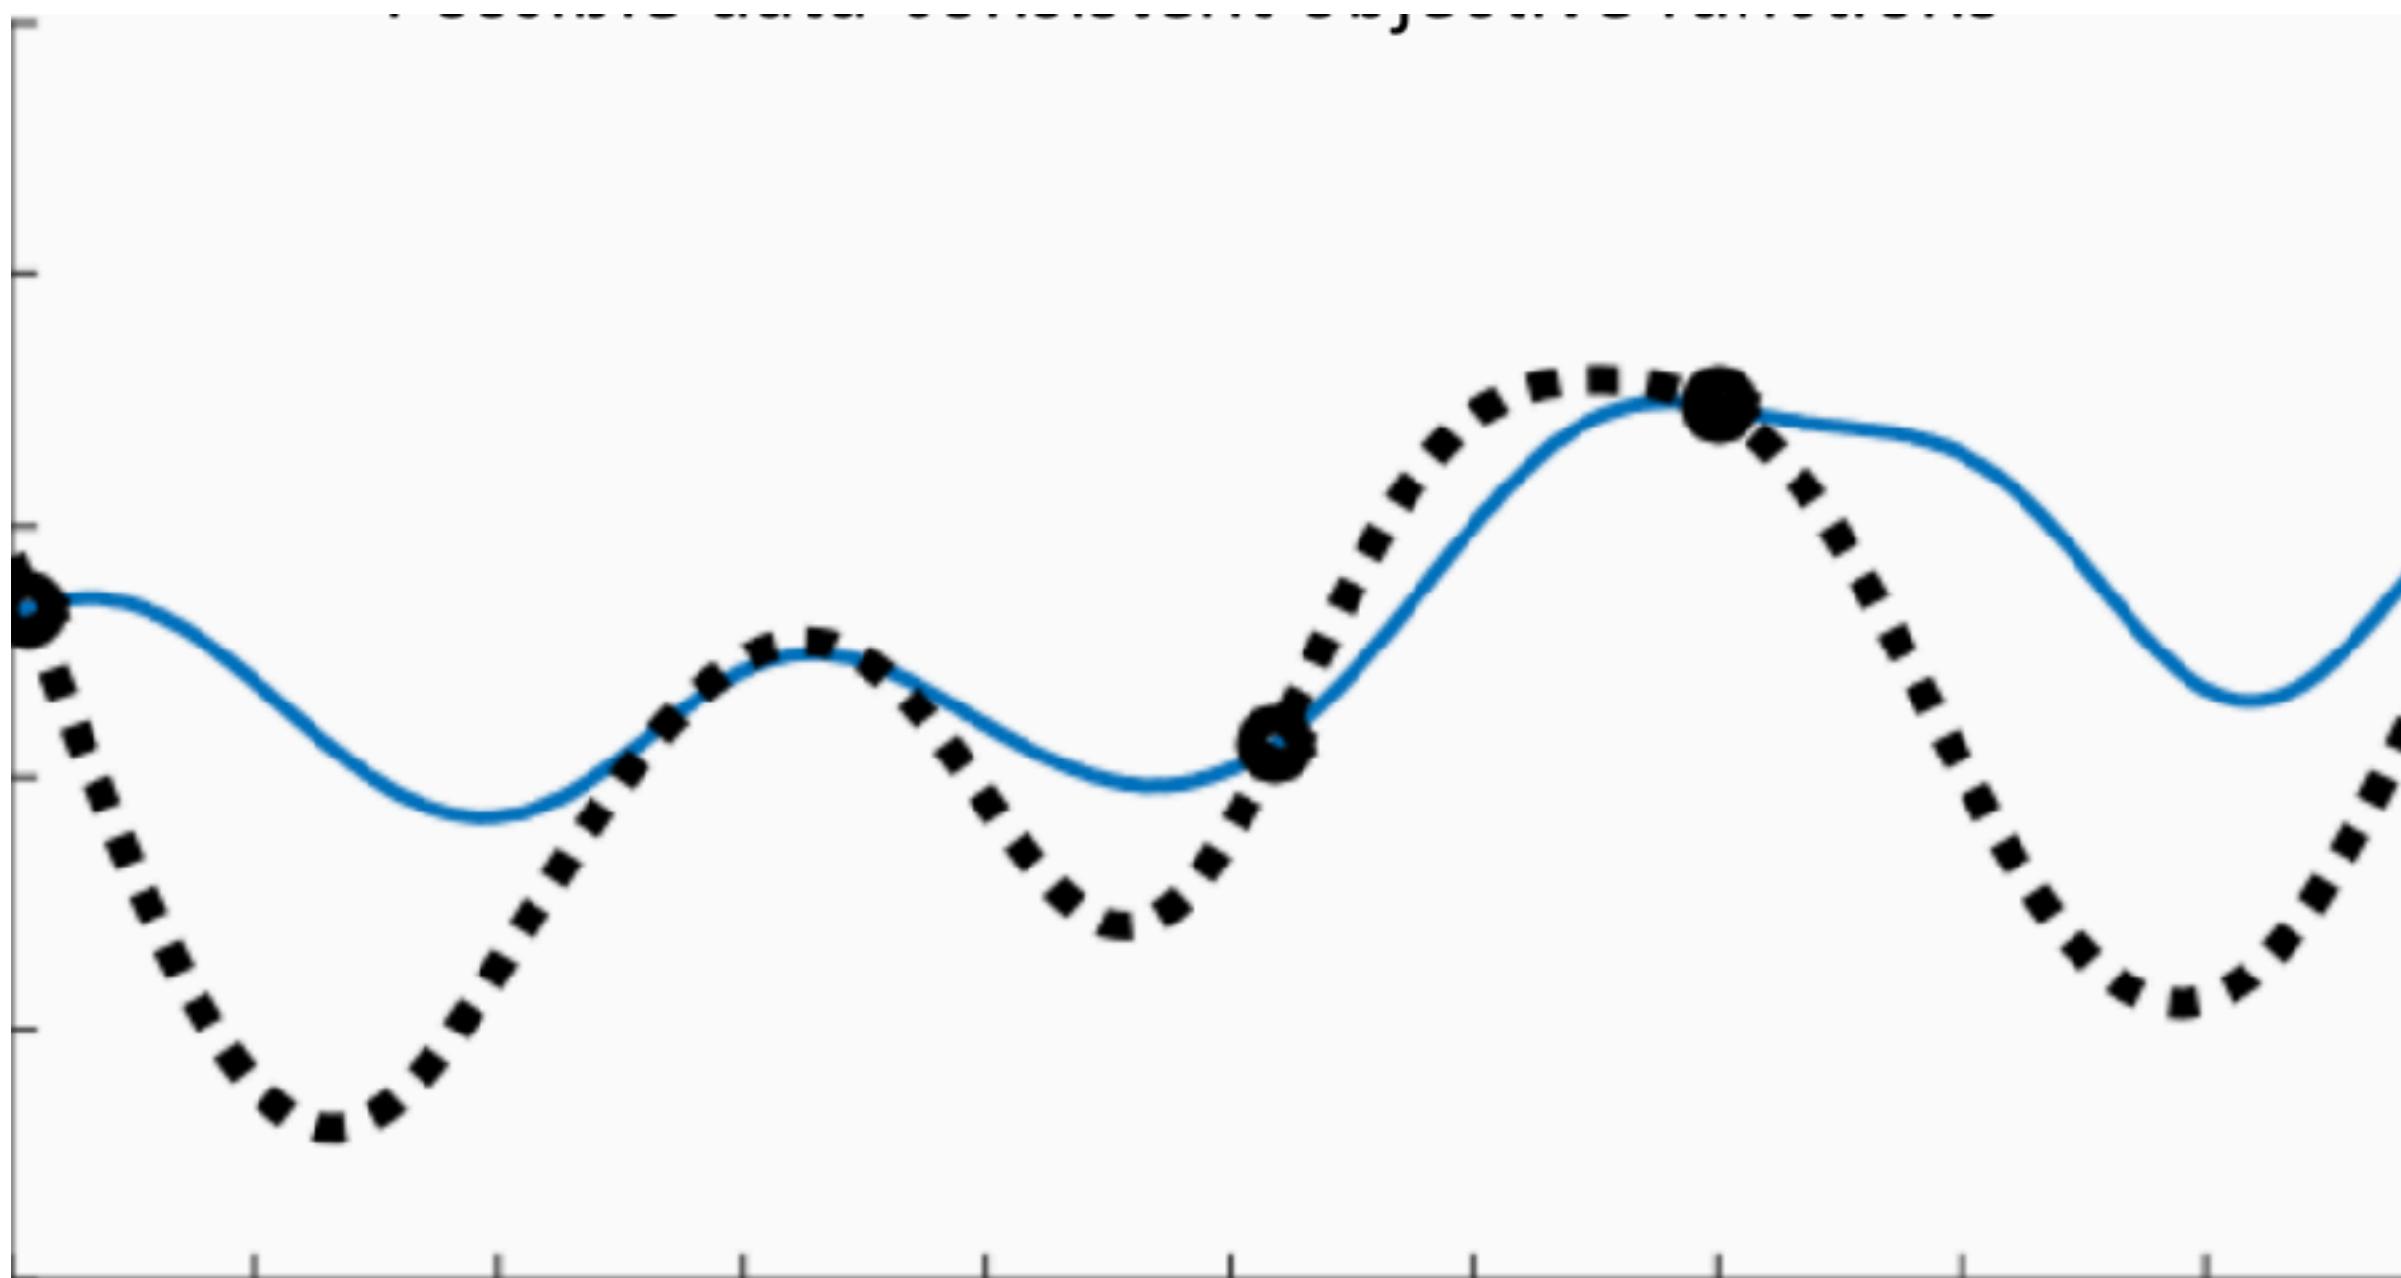
- Two different views
 - a vector of possibly uncountably many Gaussian variables with given mean and a joint covariate distribution
 - a Gaussian distribution over functions



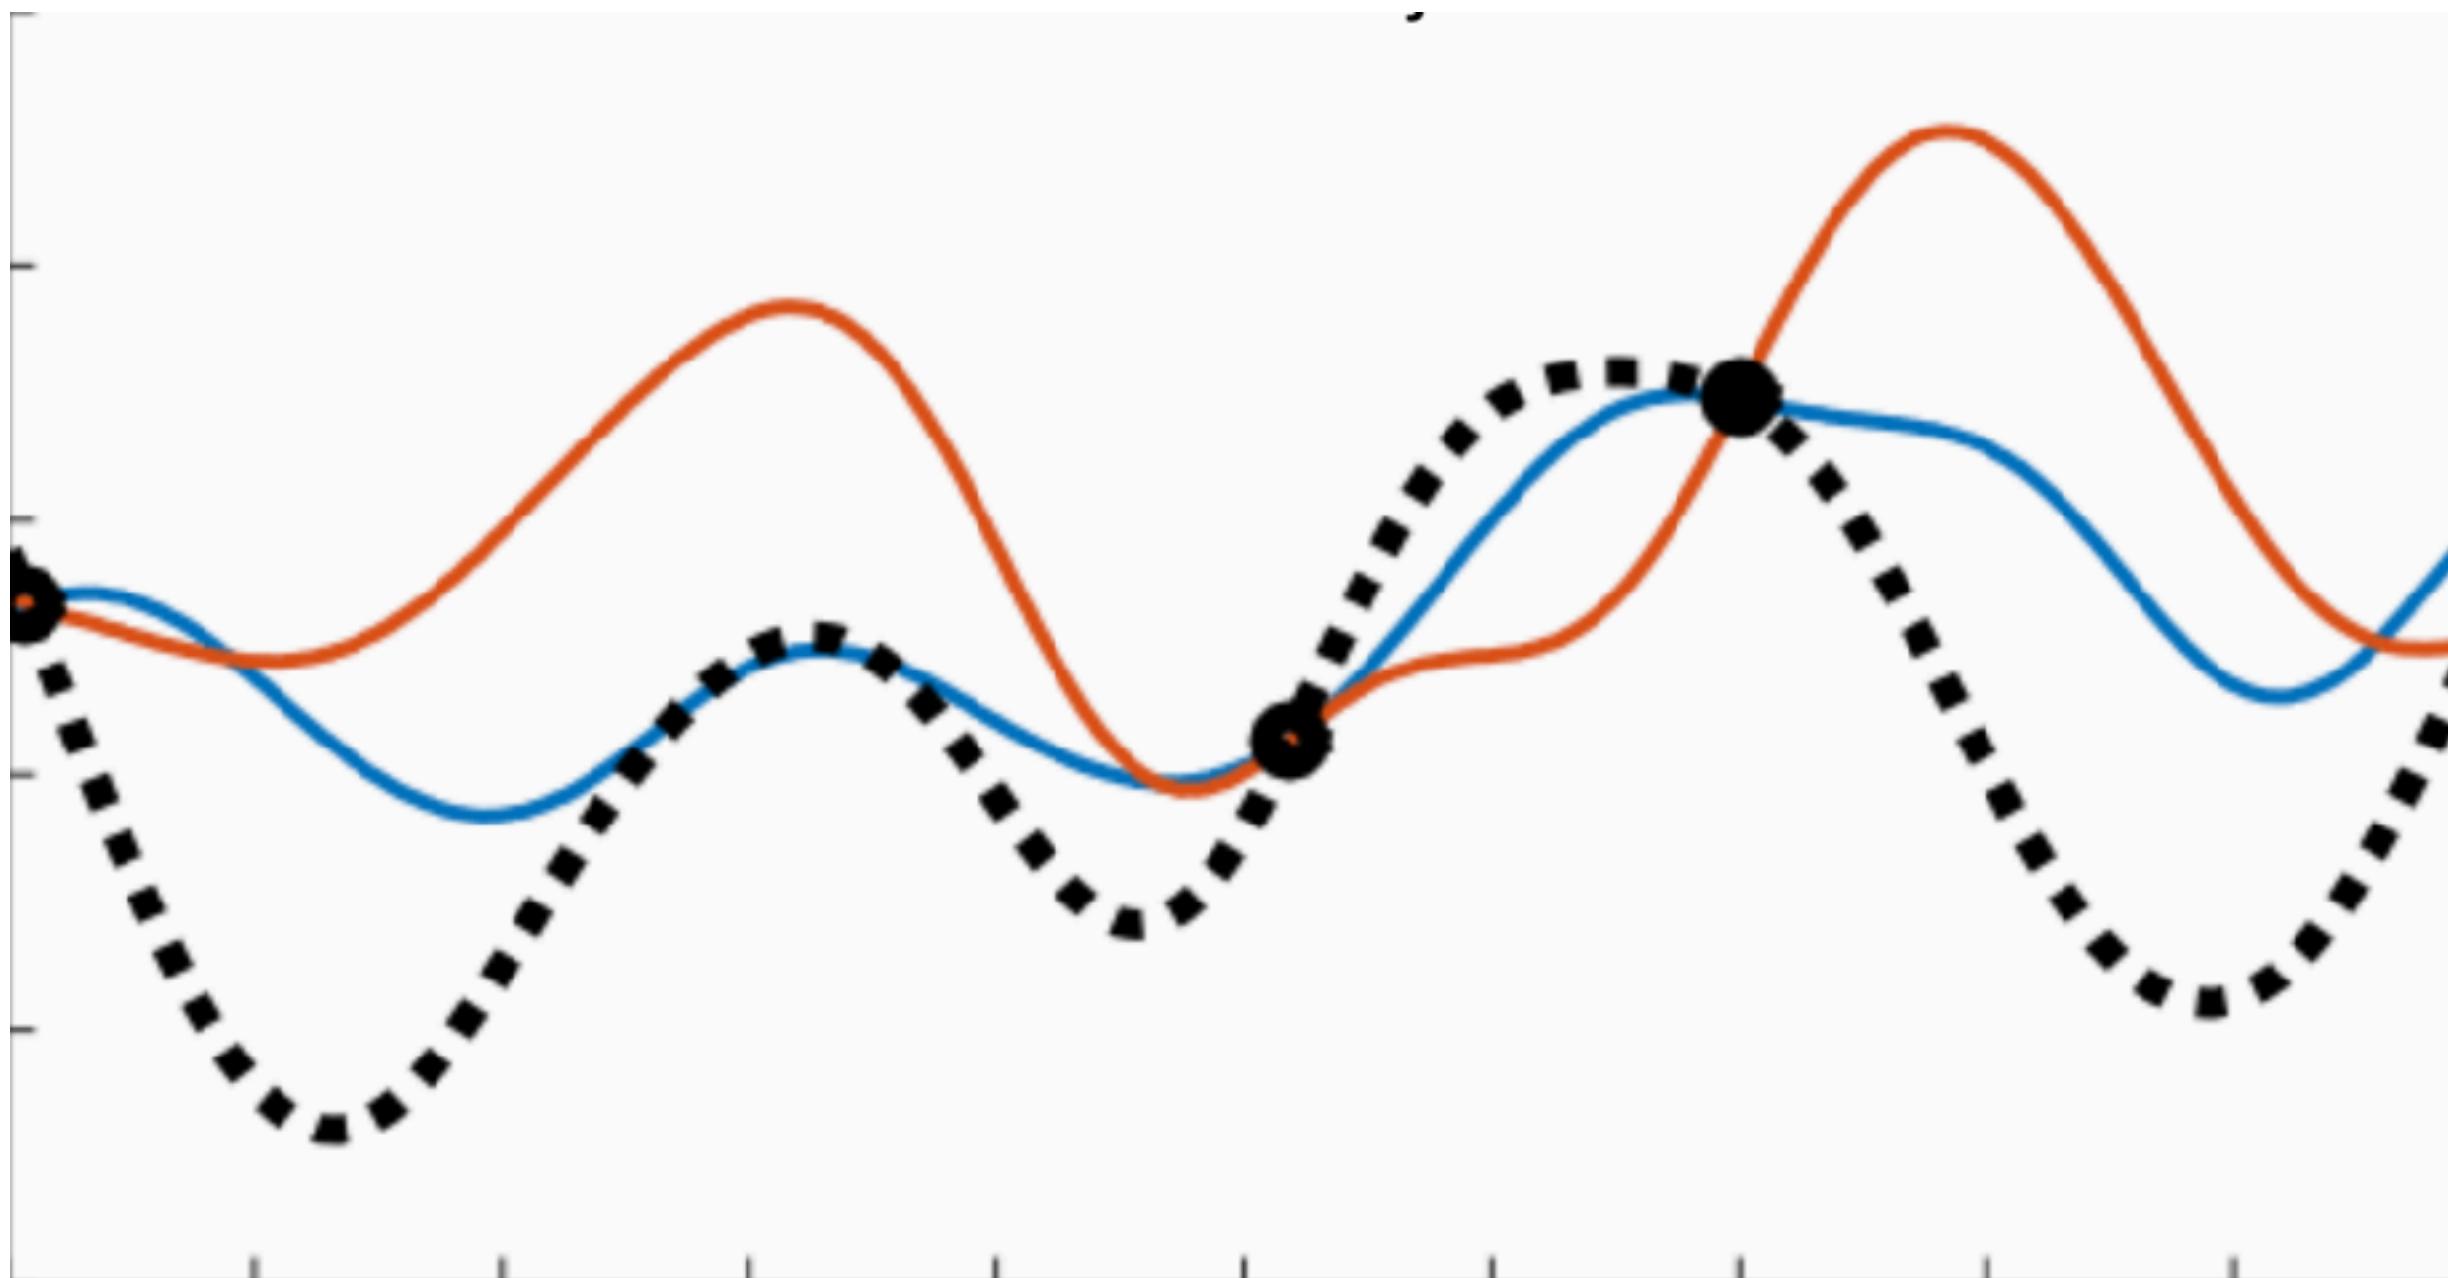
Gaussian Process



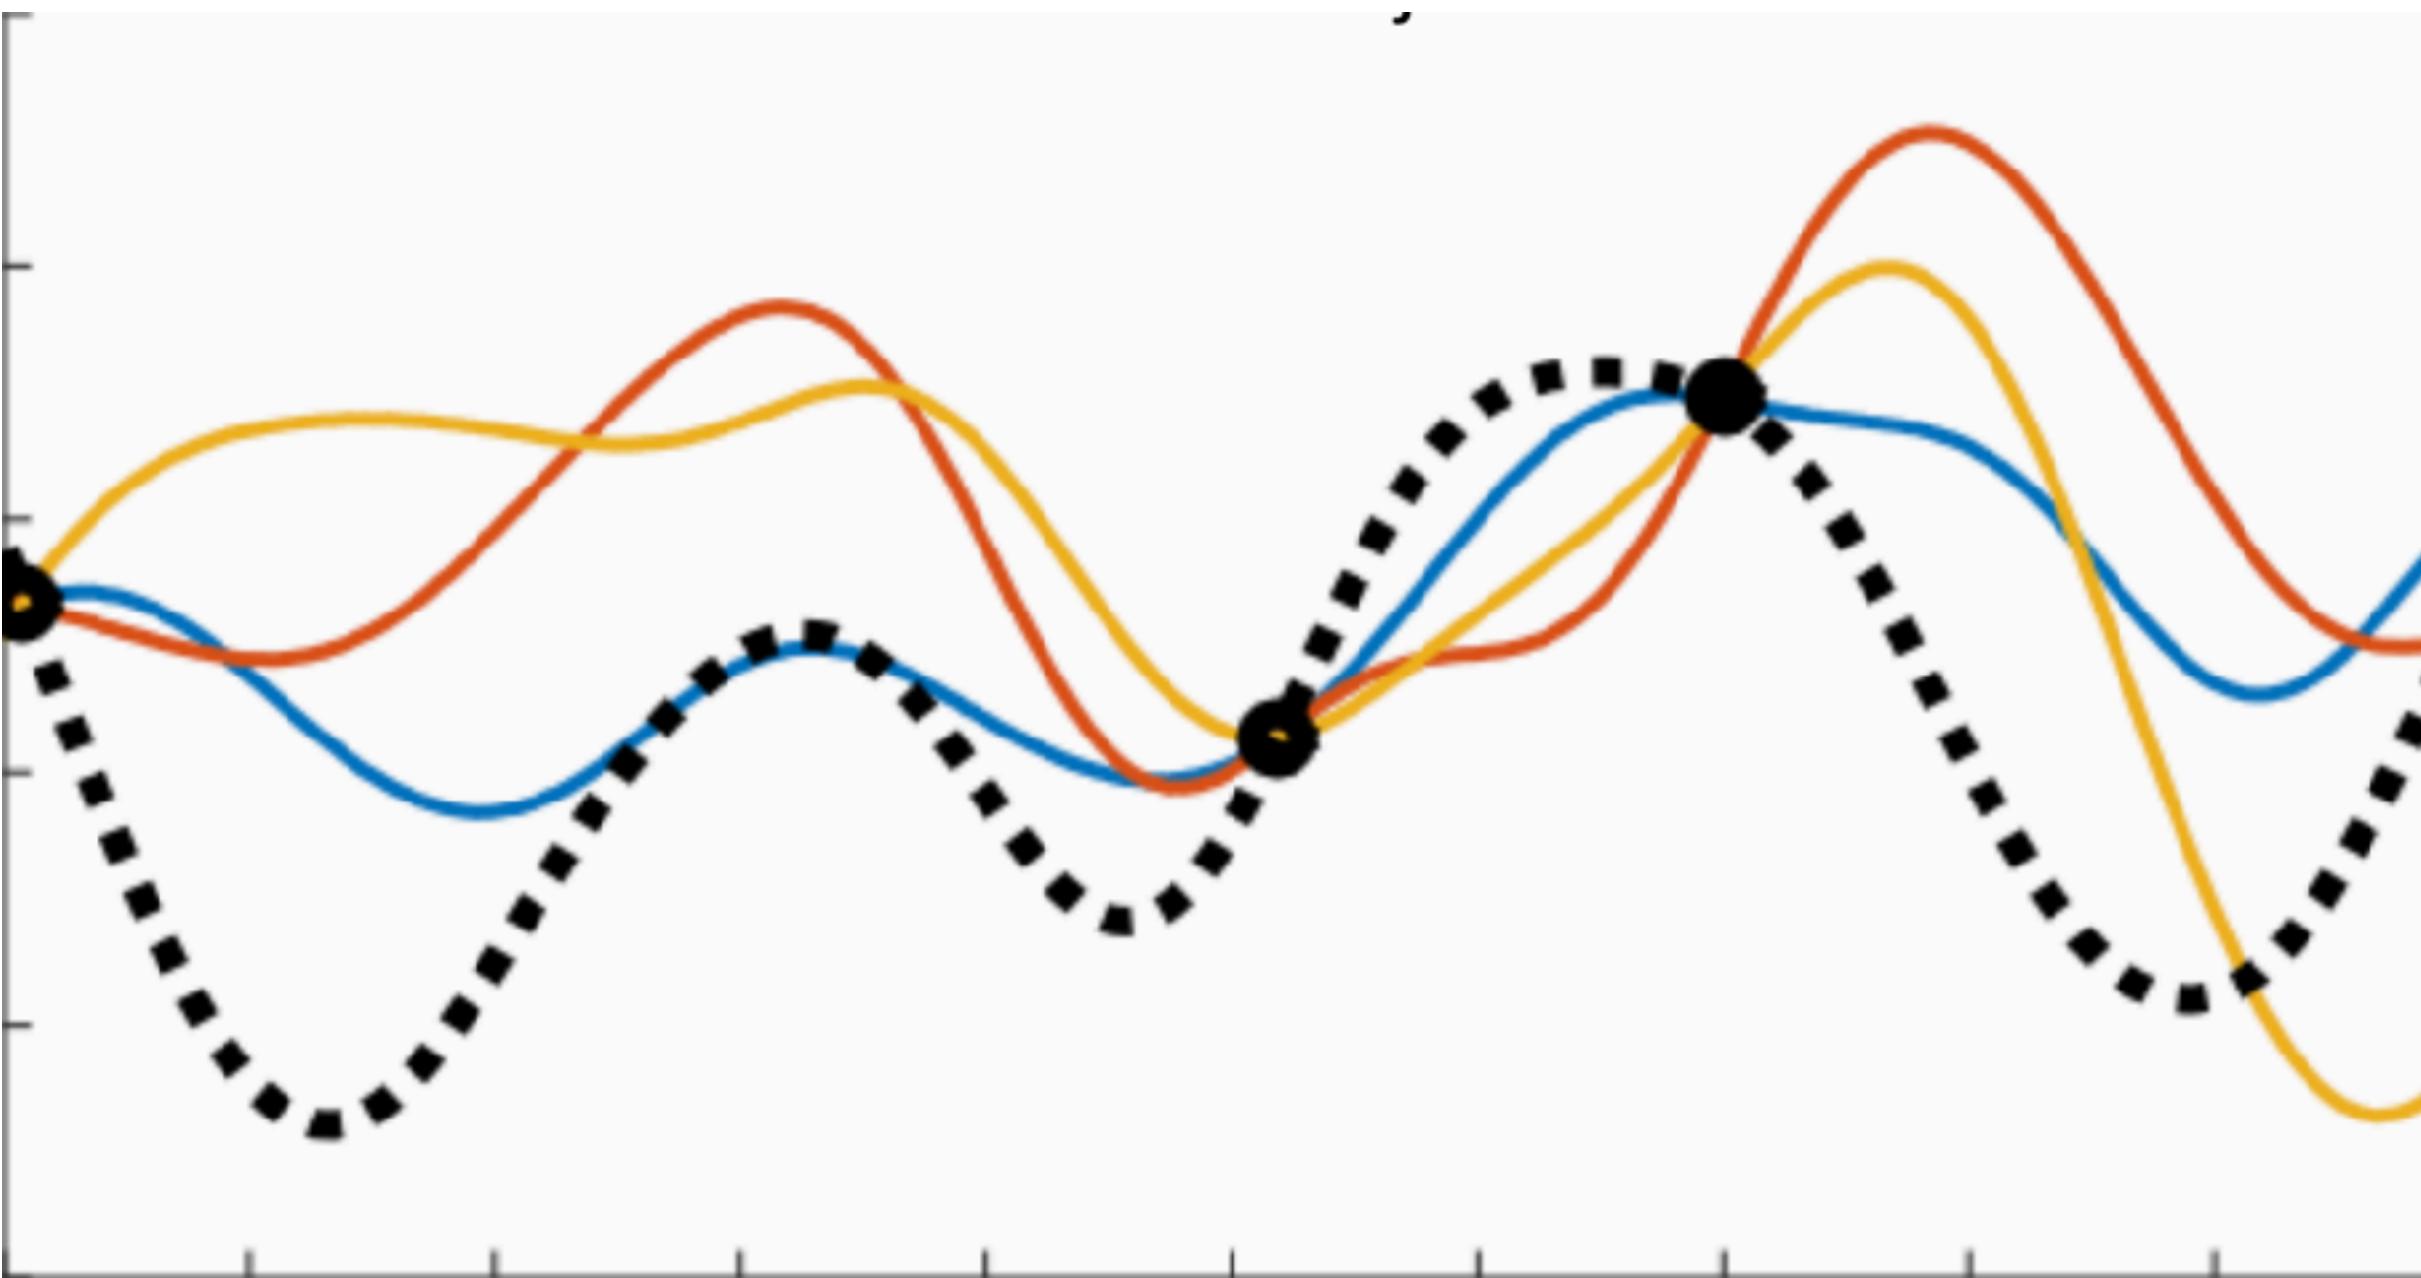
Gaussian Process



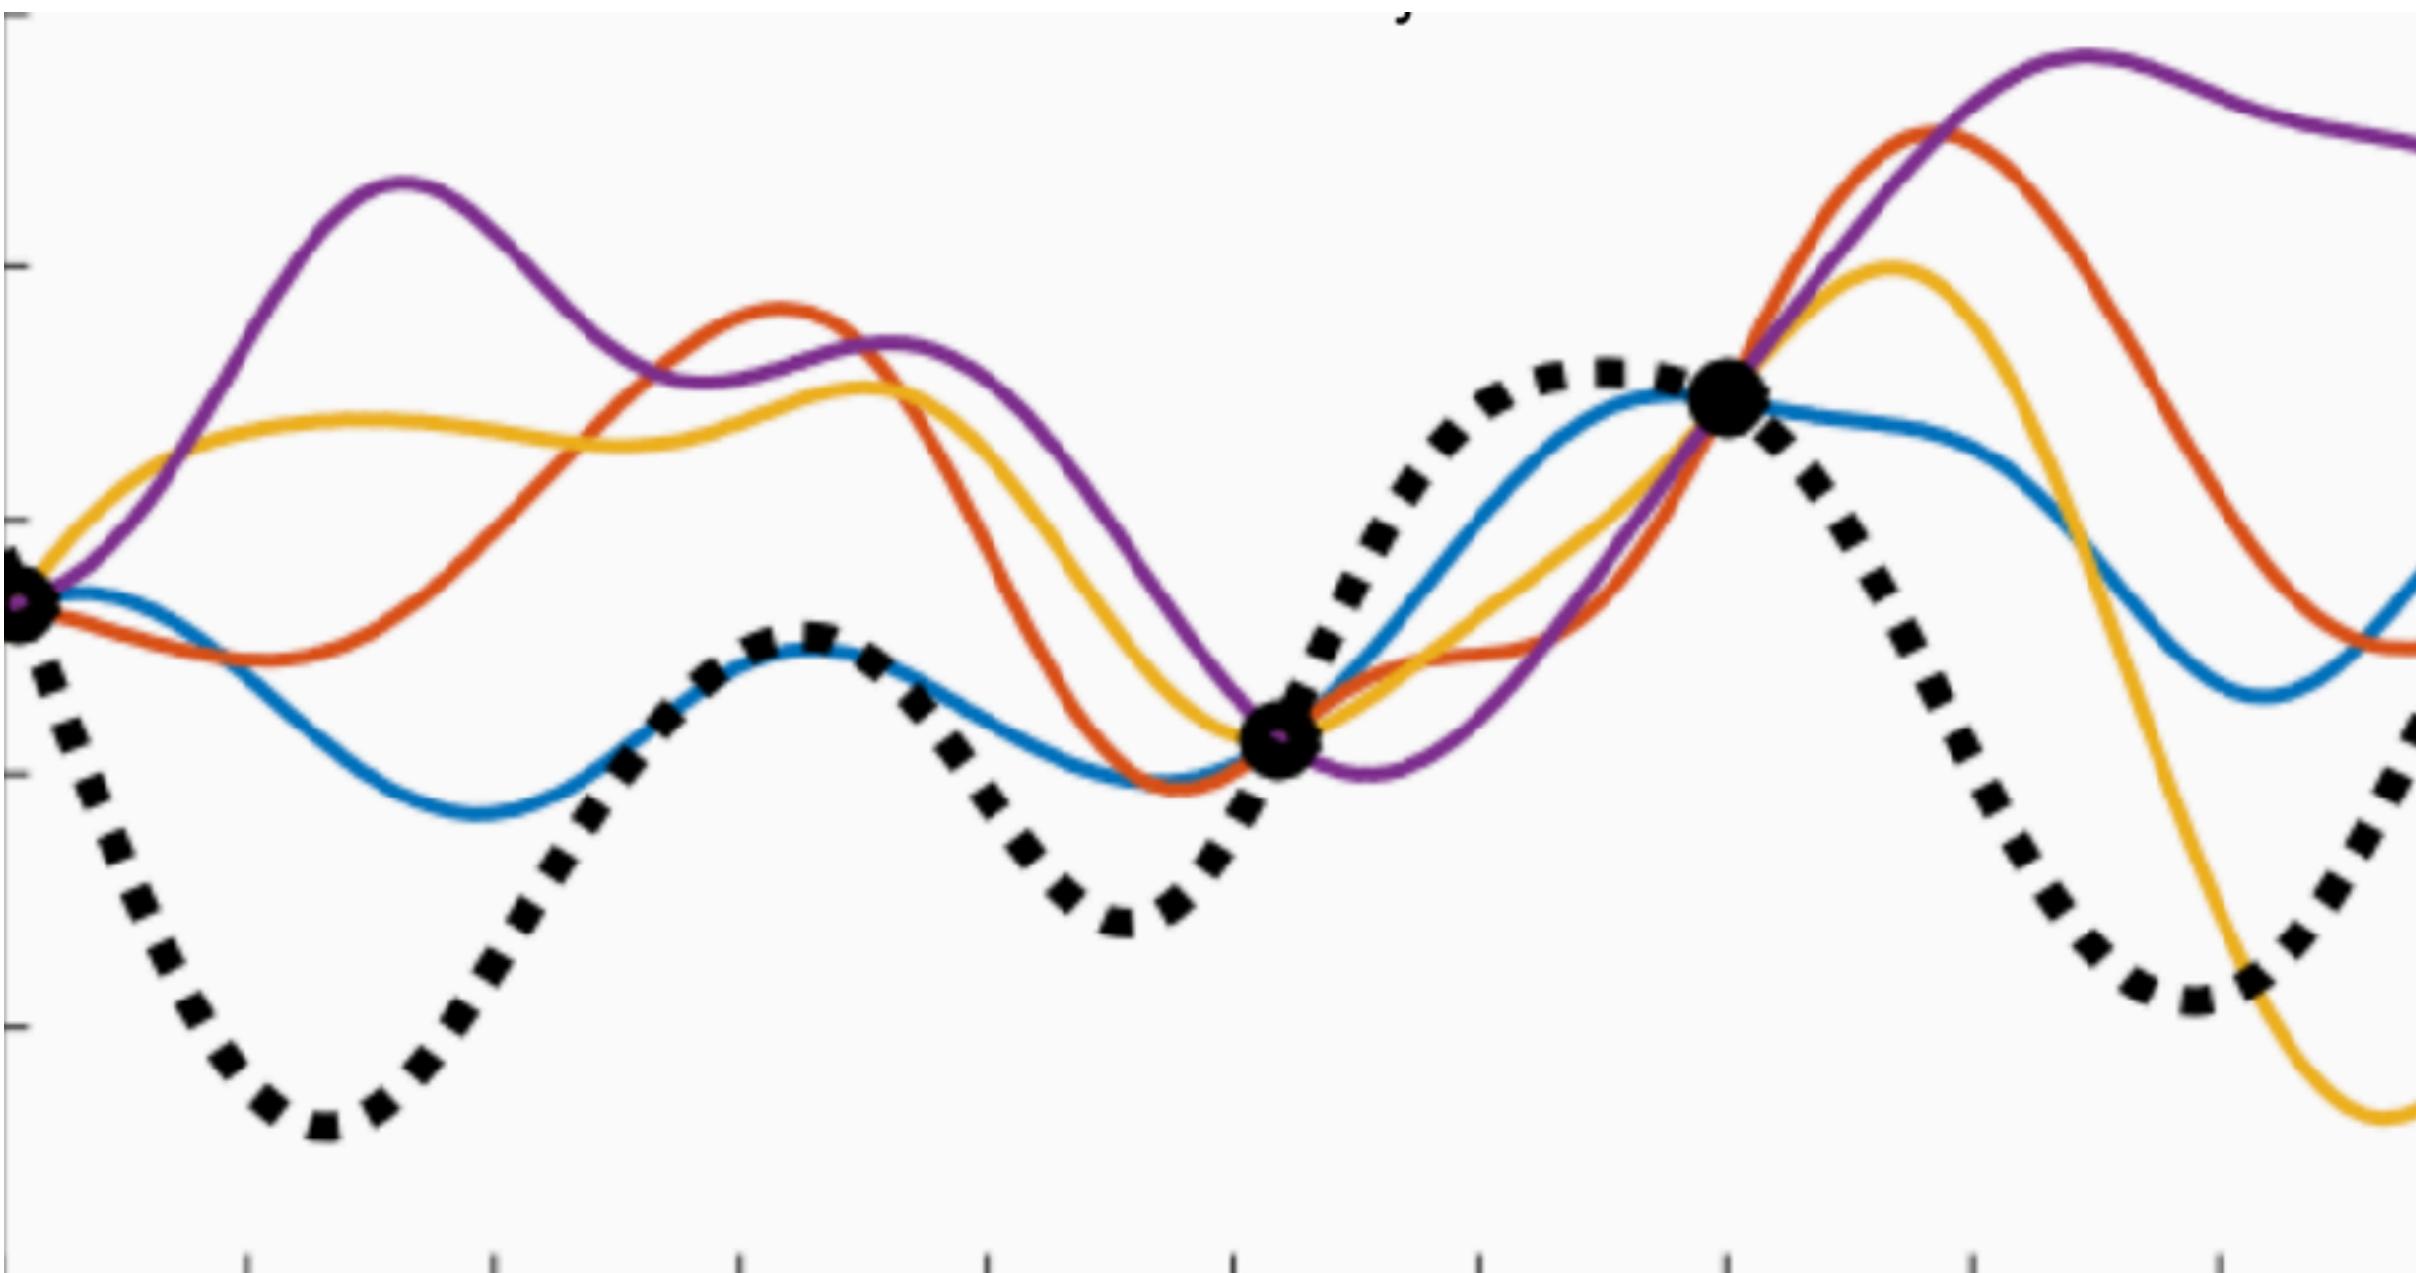
Gaussian Process



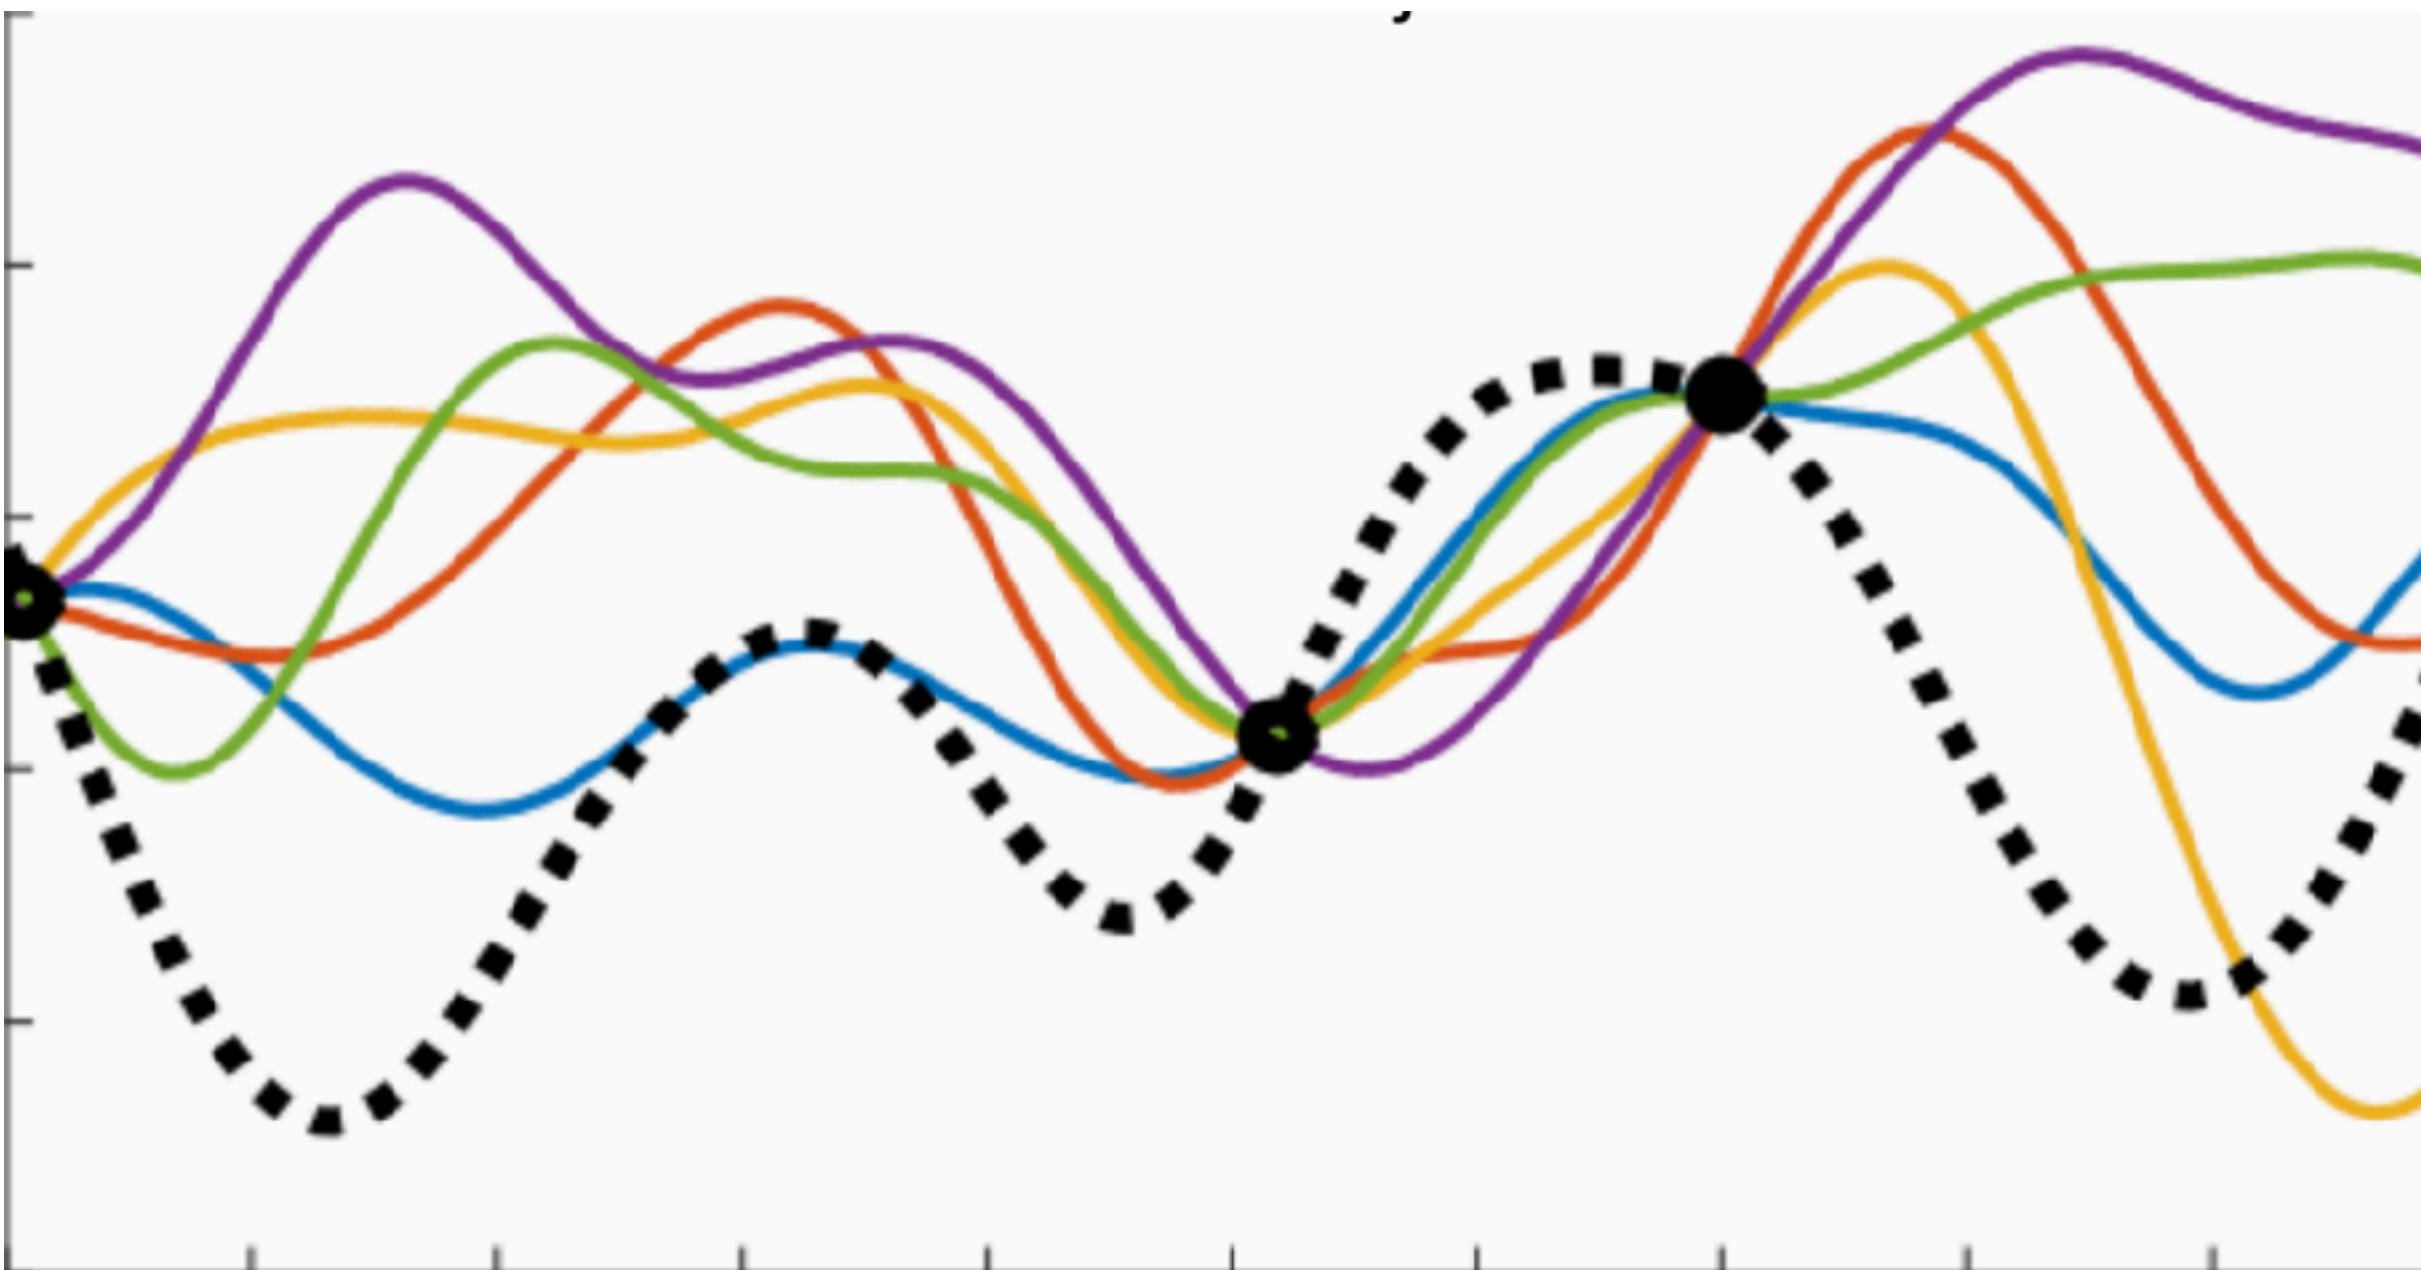
Gaussian Process



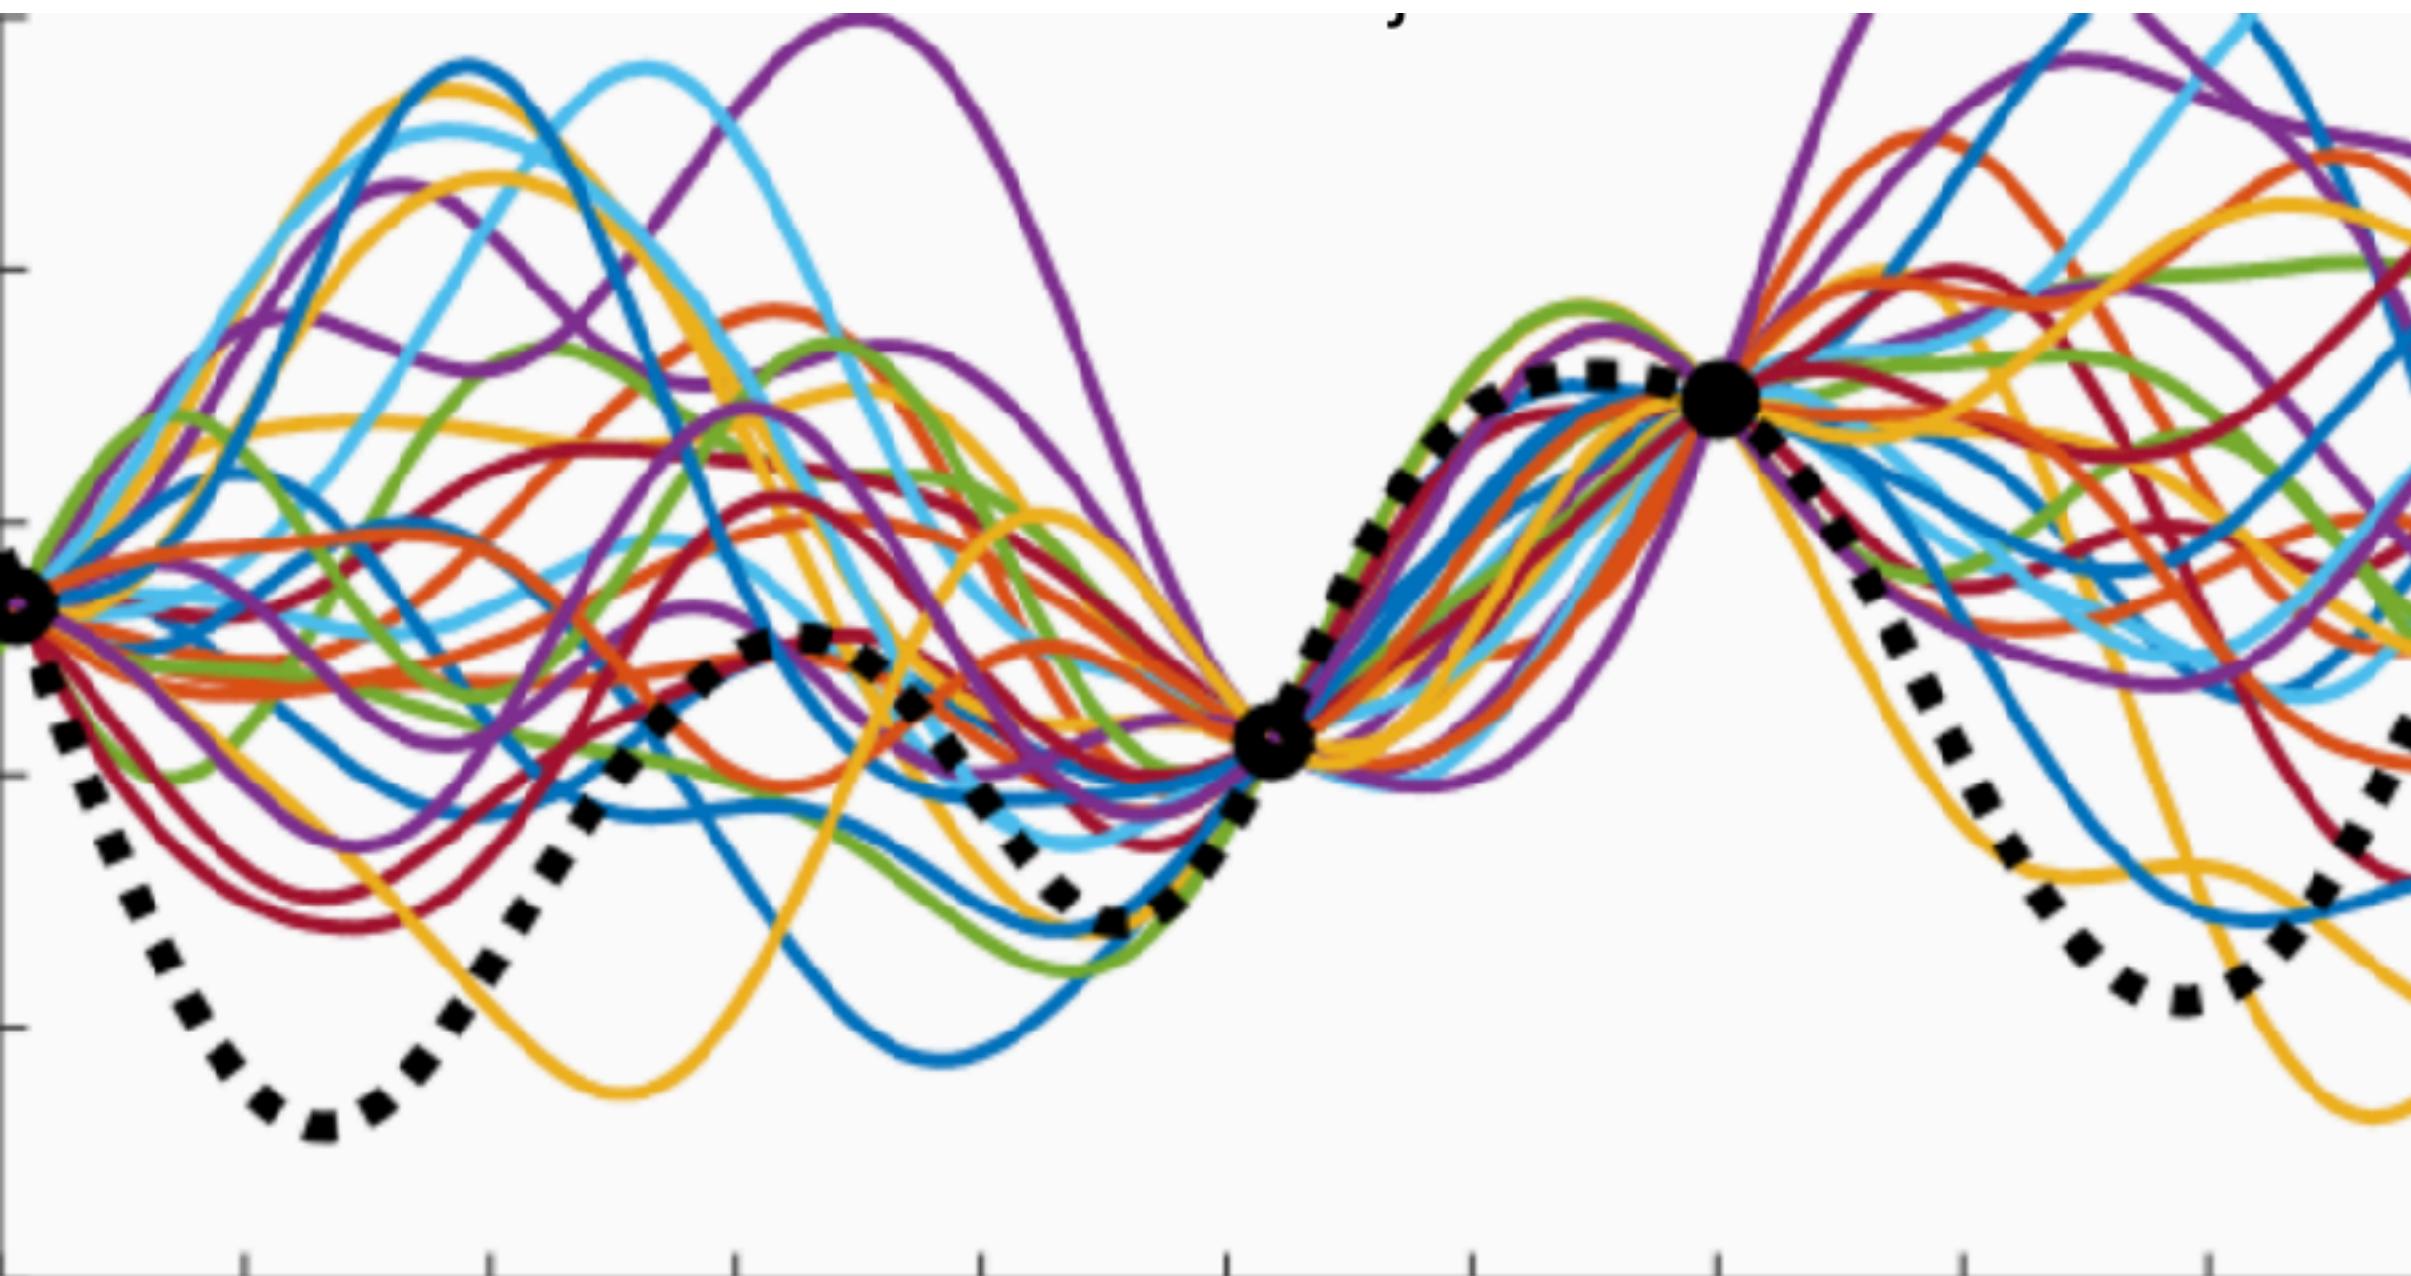
Gaussian Process



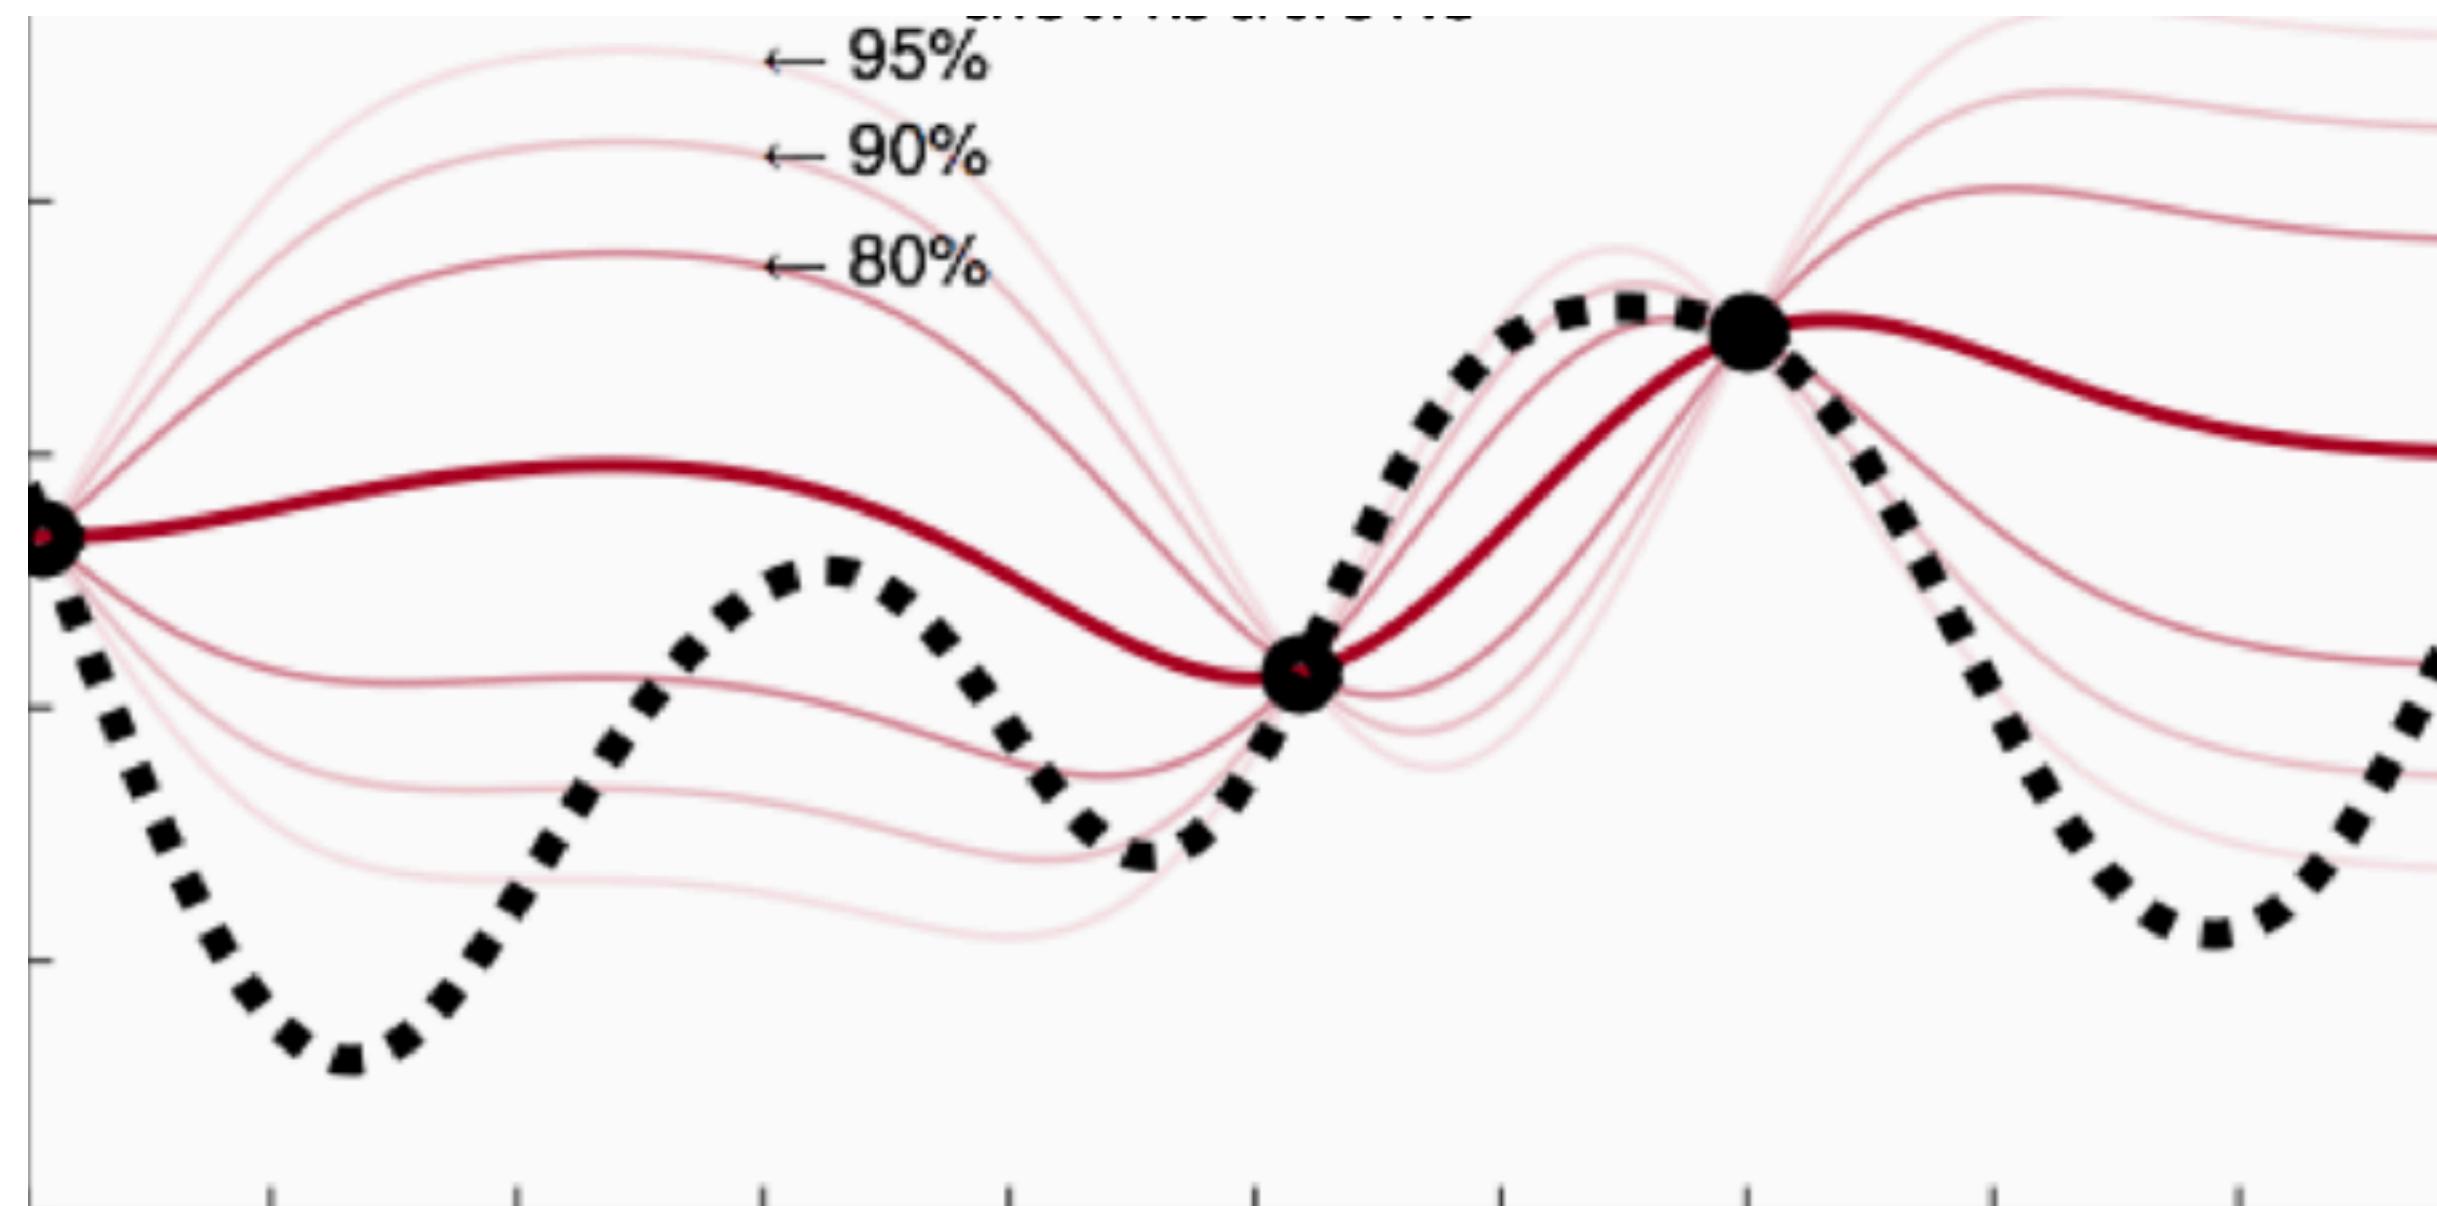
Gaussian Process



Gaussian Process



Gaussian Process



GP
n

$$\mu_n(x) = \kappa^T(\mathbf{K} + \sigma_{noise}^2 \mathbf{I})^{-1} Y$$

$$\kappa_n(x, x') = \kappa(x, x') - \kappa^T(\mathbf{K} + \sigma_{noise}^2 \mathbf{I})^{-1} \kappa'$$



Covariance Kernel Function

- Squared exponentials (SE)
- “n/2” Matern kernels



Covariance Kernel Functions

- Squared Exponential:

$$k_i(x_i - x'_i) = e^{-\left(\frac{x_i - x'_i}{\theta_i}\right)^2}$$

- Matern Function $\nu = \frac{3}{2}$:

$$k_i(x_i - x'_i) = \left(1 + \sqrt{3} \left|\frac{x_i - x'_i}{\theta_i}\right|\right) e^{-\sqrt{3} \left|\frac{x_i - x'_i}{\theta_i}\right|}$$

- Matern Function $\nu = \frac{5}{2}$:

$$k_i(x_i - x'_i) = \left(1 + \sqrt{5} \left|\frac{x_i - x'_i}{\theta_i}\right| + \frac{5}{3} \left|\frac{x_i - x'_i}{\theta_i}\right|^2\right) e^{-\sqrt{5} \left|\frac{x_i - x'_i}{\theta_i}\right|}$$



Covariance Kernel Functions

- Squared Exponential:

$$k_i(x_i - x'_i) = e^{-\left(\frac{x_i - x'_i}{\theta_i}\right)^2}$$

- Matern Function $\nu = \frac{3}{2}$:

$$k_i(x_i - x'_i) = \left(1 + \sqrt{3} \left|\frac{x_i - x'_i}{\theta_i}\right| + \frac{3}{4} \left(\frac{x_i - x'_i}{\theta_i}\right)^2\right) e^{-\sqrt{3} \left|\frac{x_i - x'_i}{\theta_i}\right|}$$

- Matern Function $\nu = \frac{5}{2}$:

$$k_i(x_i - x'_i) = \left(1 + \sqrt{5} \left|\frac{x_i - x'_i}{\theta_i}\right| + \frac{5}{3} \left(\frac{x_i - x'_i}{\theta_i}\right)^2 + \frac{5}{8} \left(\frac{x_i - x'_i}{\theta_i}\right)^4\right) e^{-\sqrt{5} \left|\frac{x_i - x'_i}{\theta_i}\right|}$$



Acquisition Function

$$GP_{prior} + Data_n \xrightarrow{Bayes} GP_n \xrightarrow{?} x_{n+1}$$


Acquisition Function

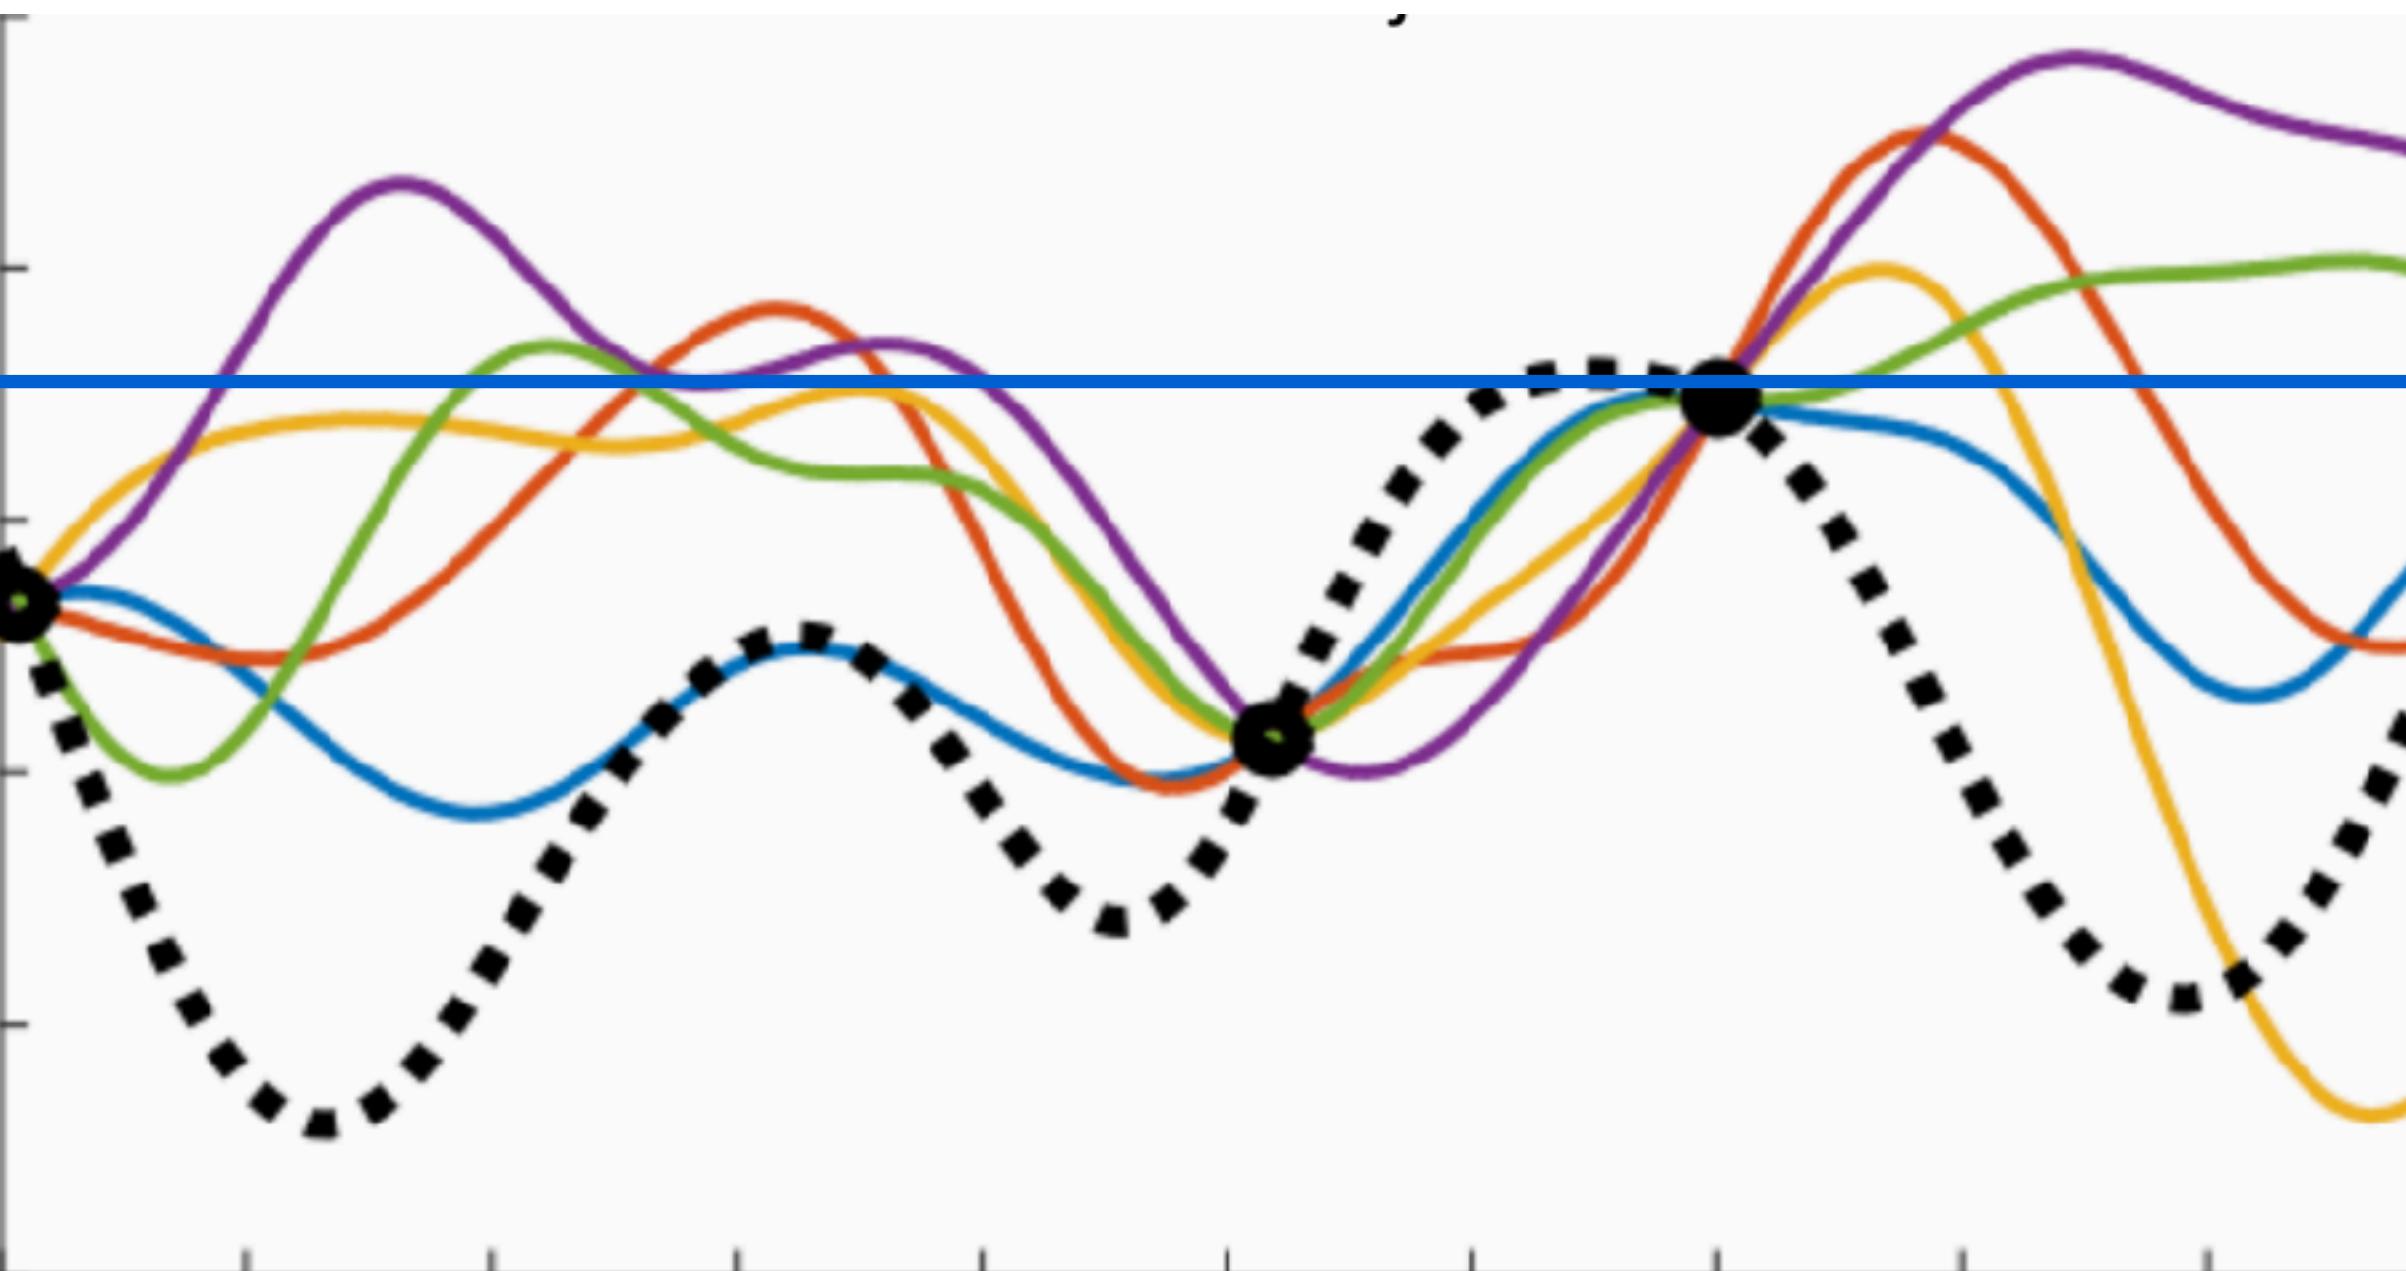
$$GP_{prior} + Data_{n+1} \xrightarrow{Bayes} GP_{n+1} \xrightarrow{?} x_{n+2}$$


Acquisition Functions

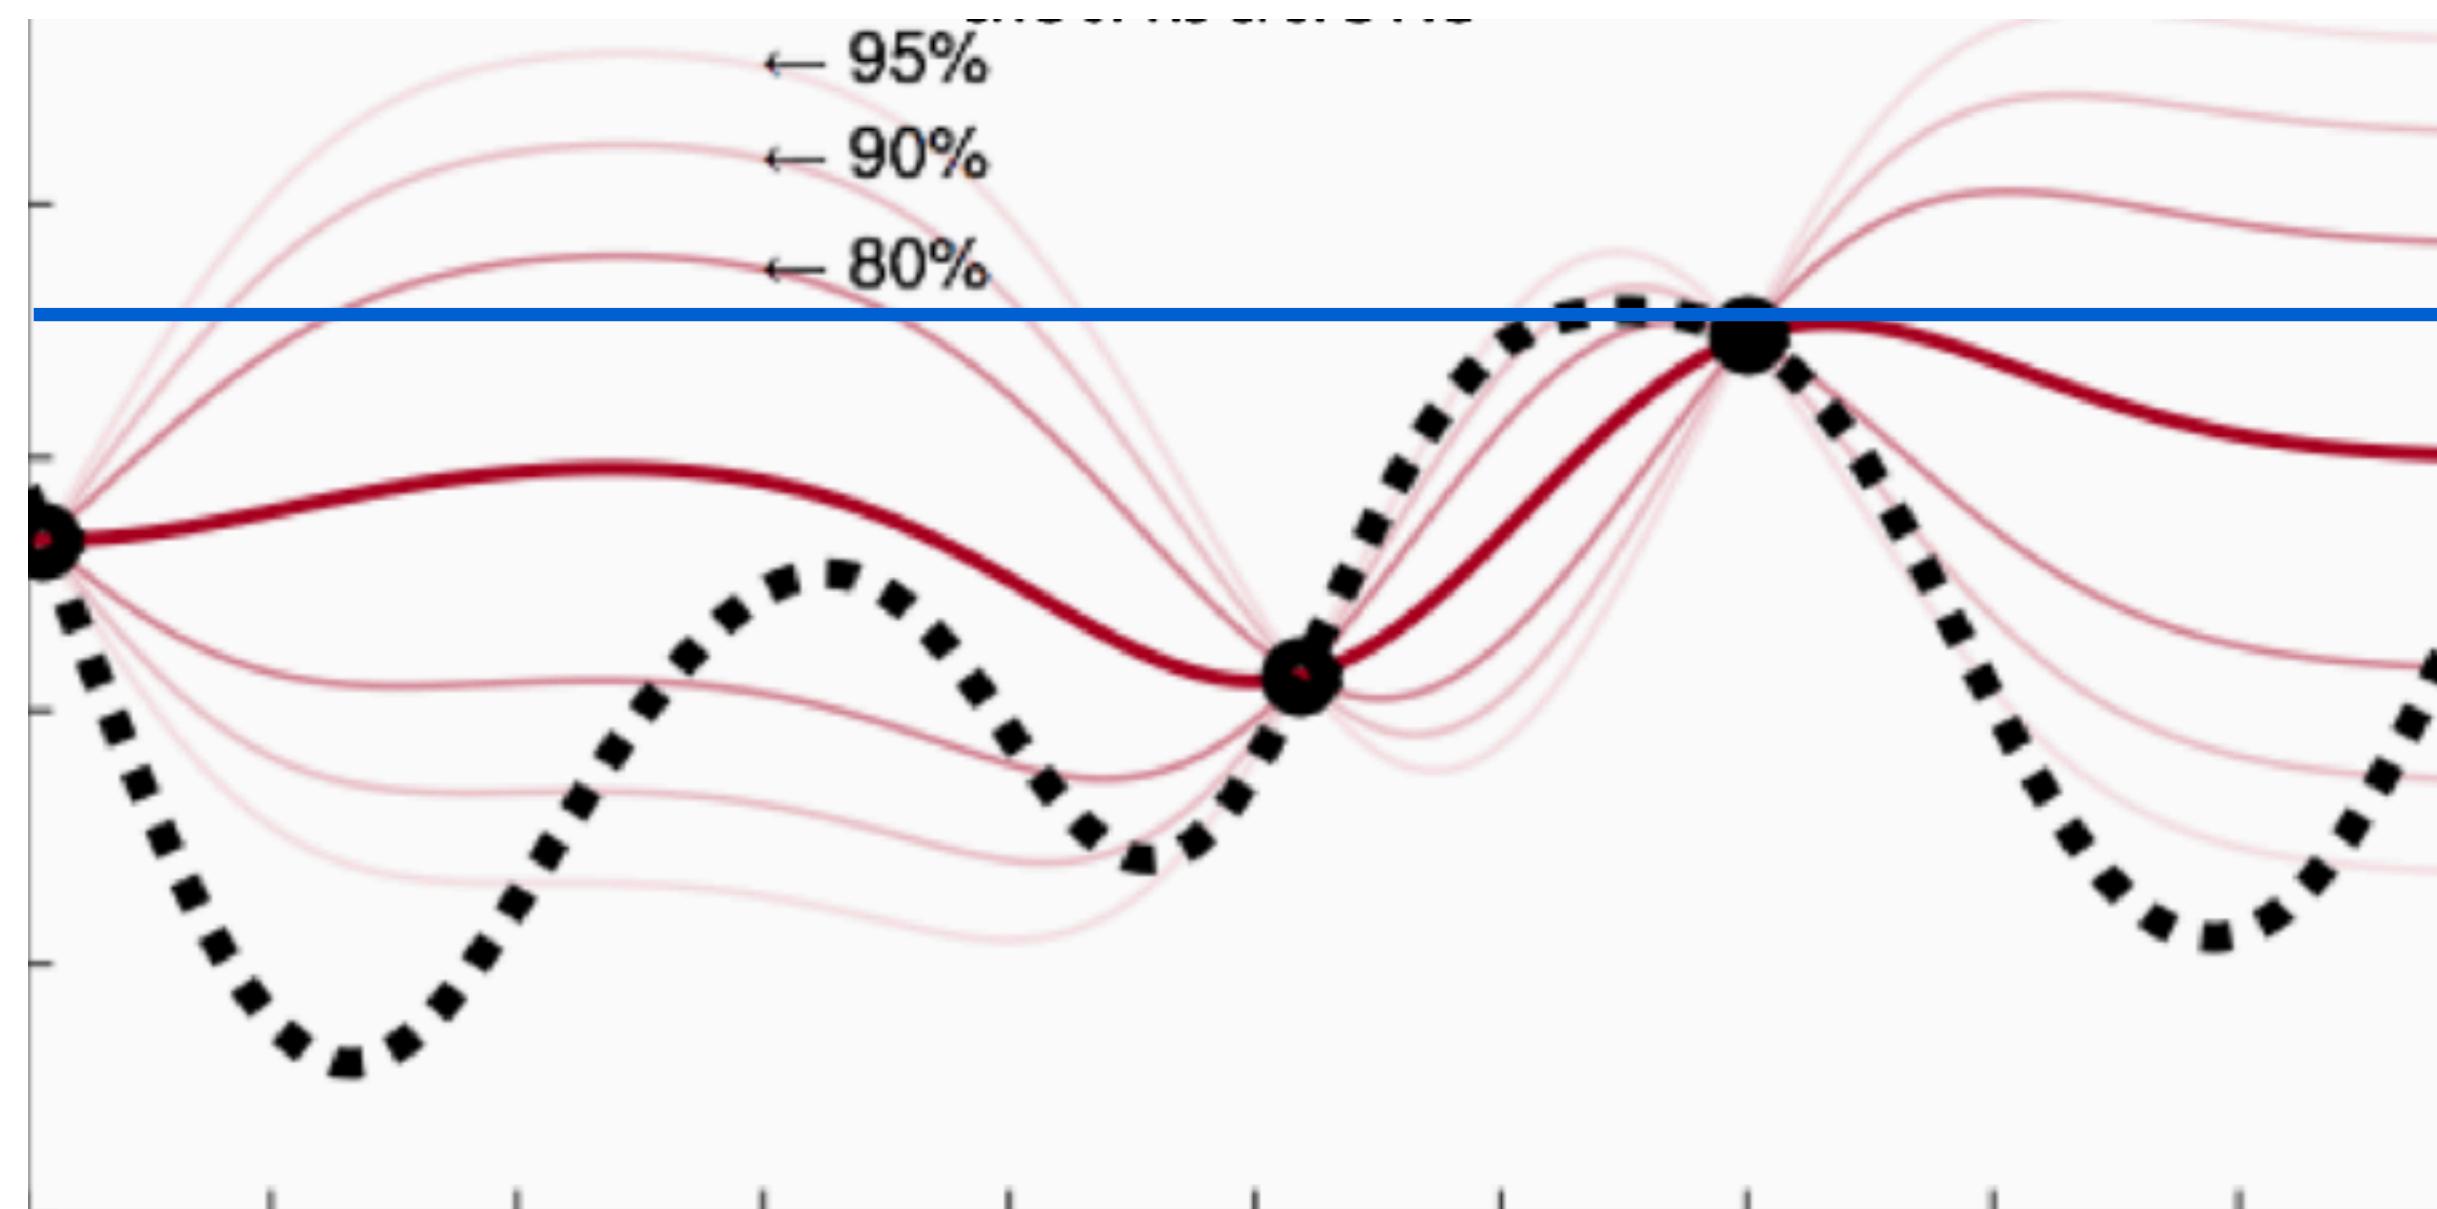
- Thompson Sampling from the posterior GP (TS)
- Probability of Improvement (PI)
- Upper Confidence Bound (UCB)
- Expected Improvement (EI)



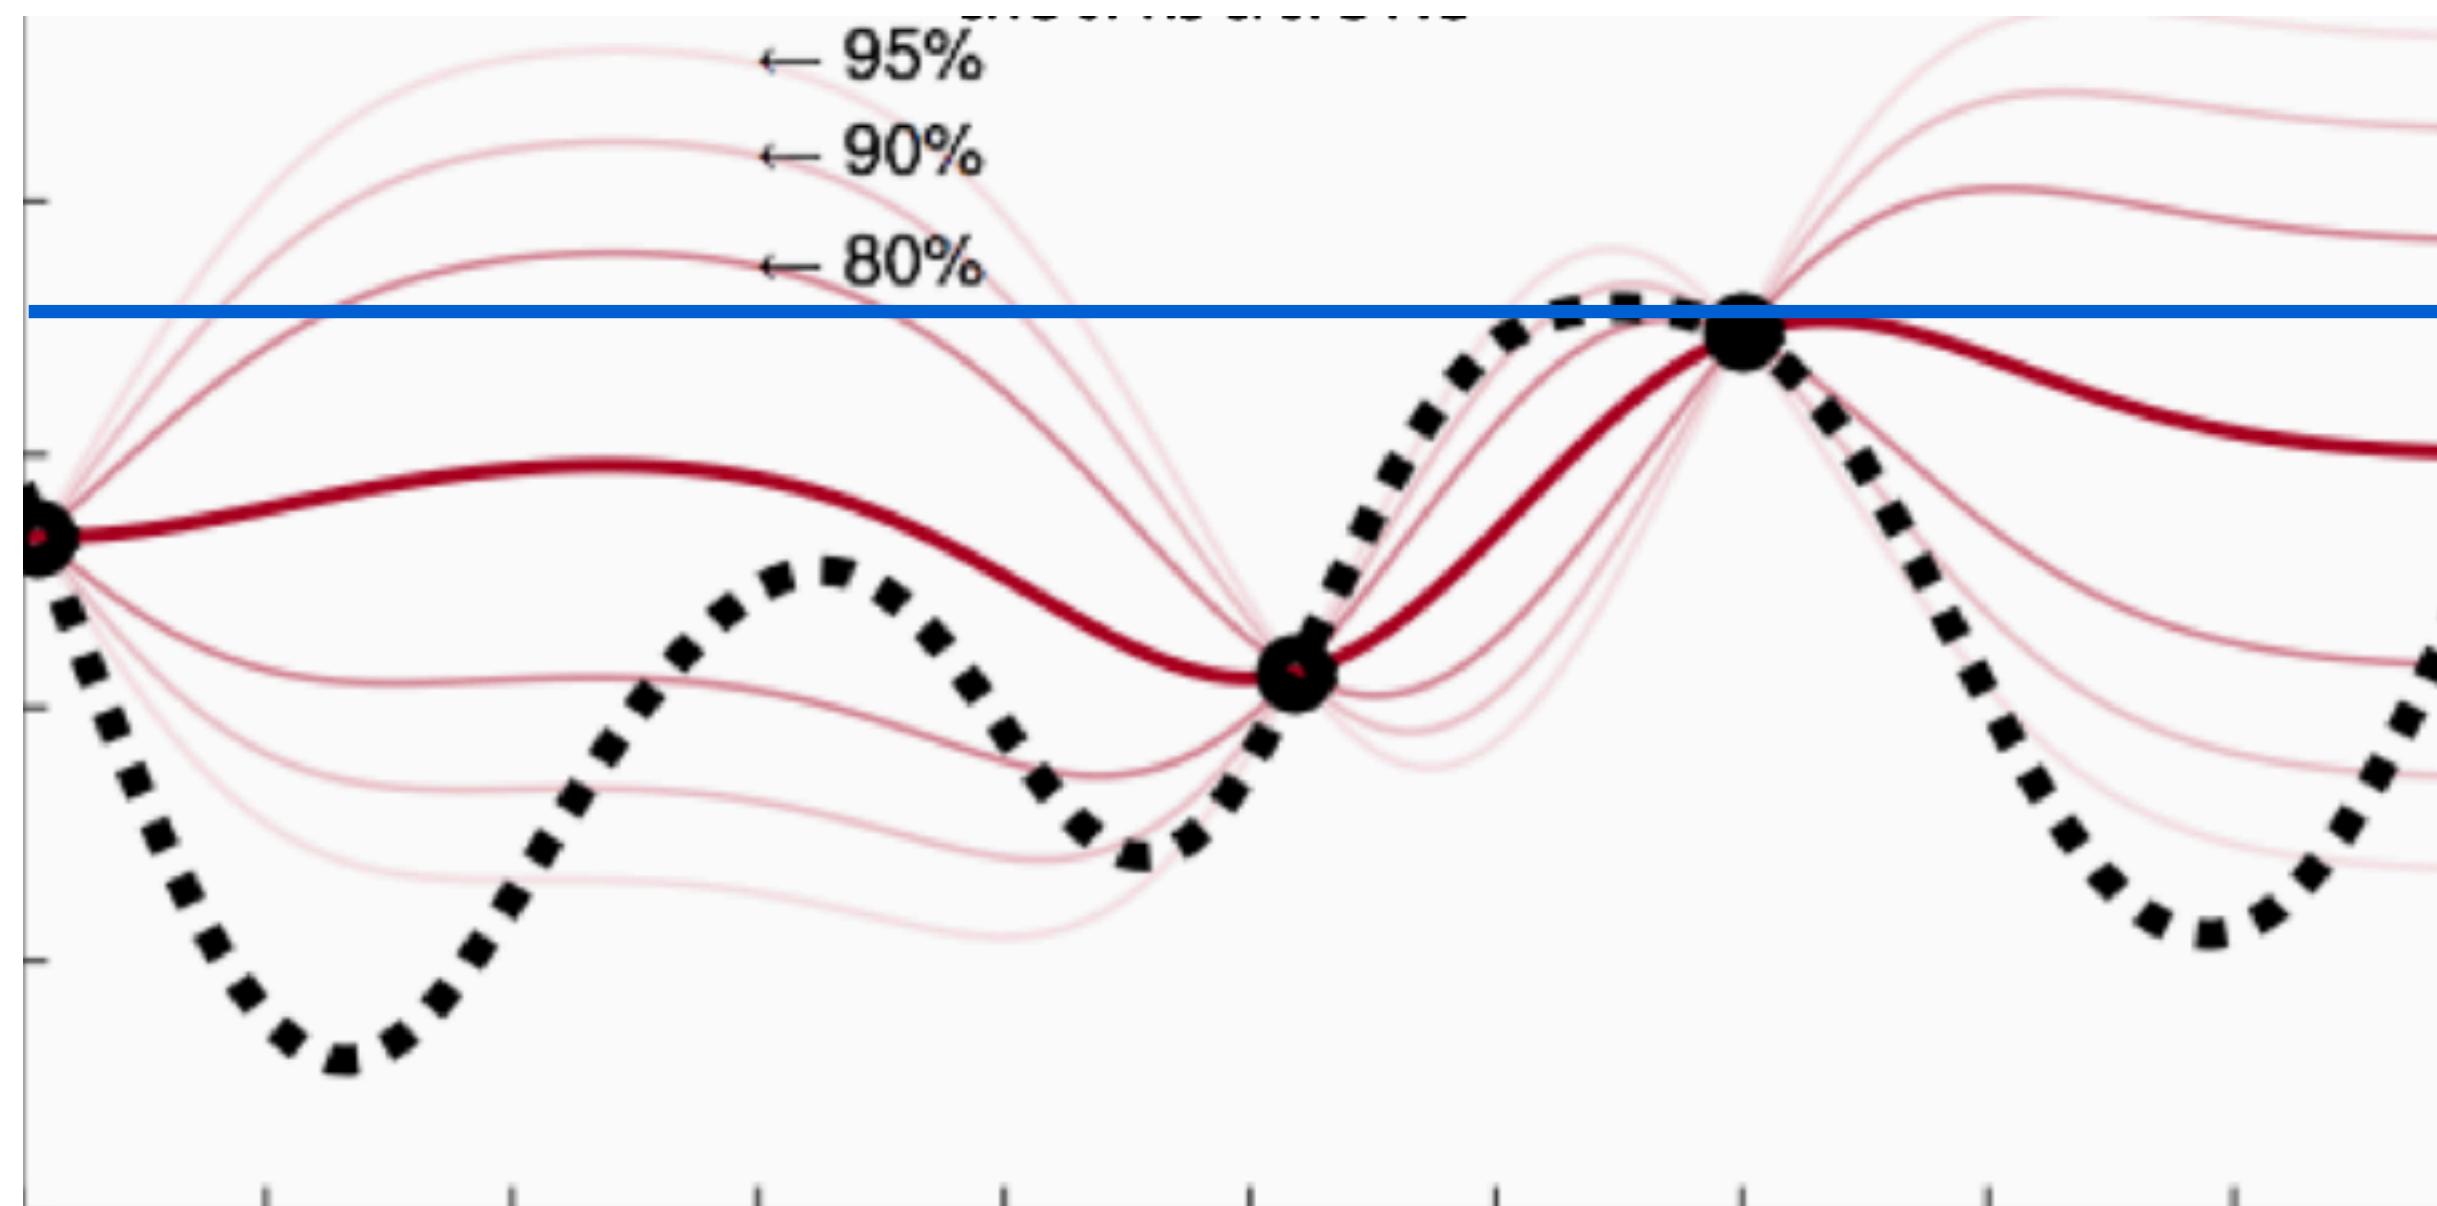
Thompson Sampling



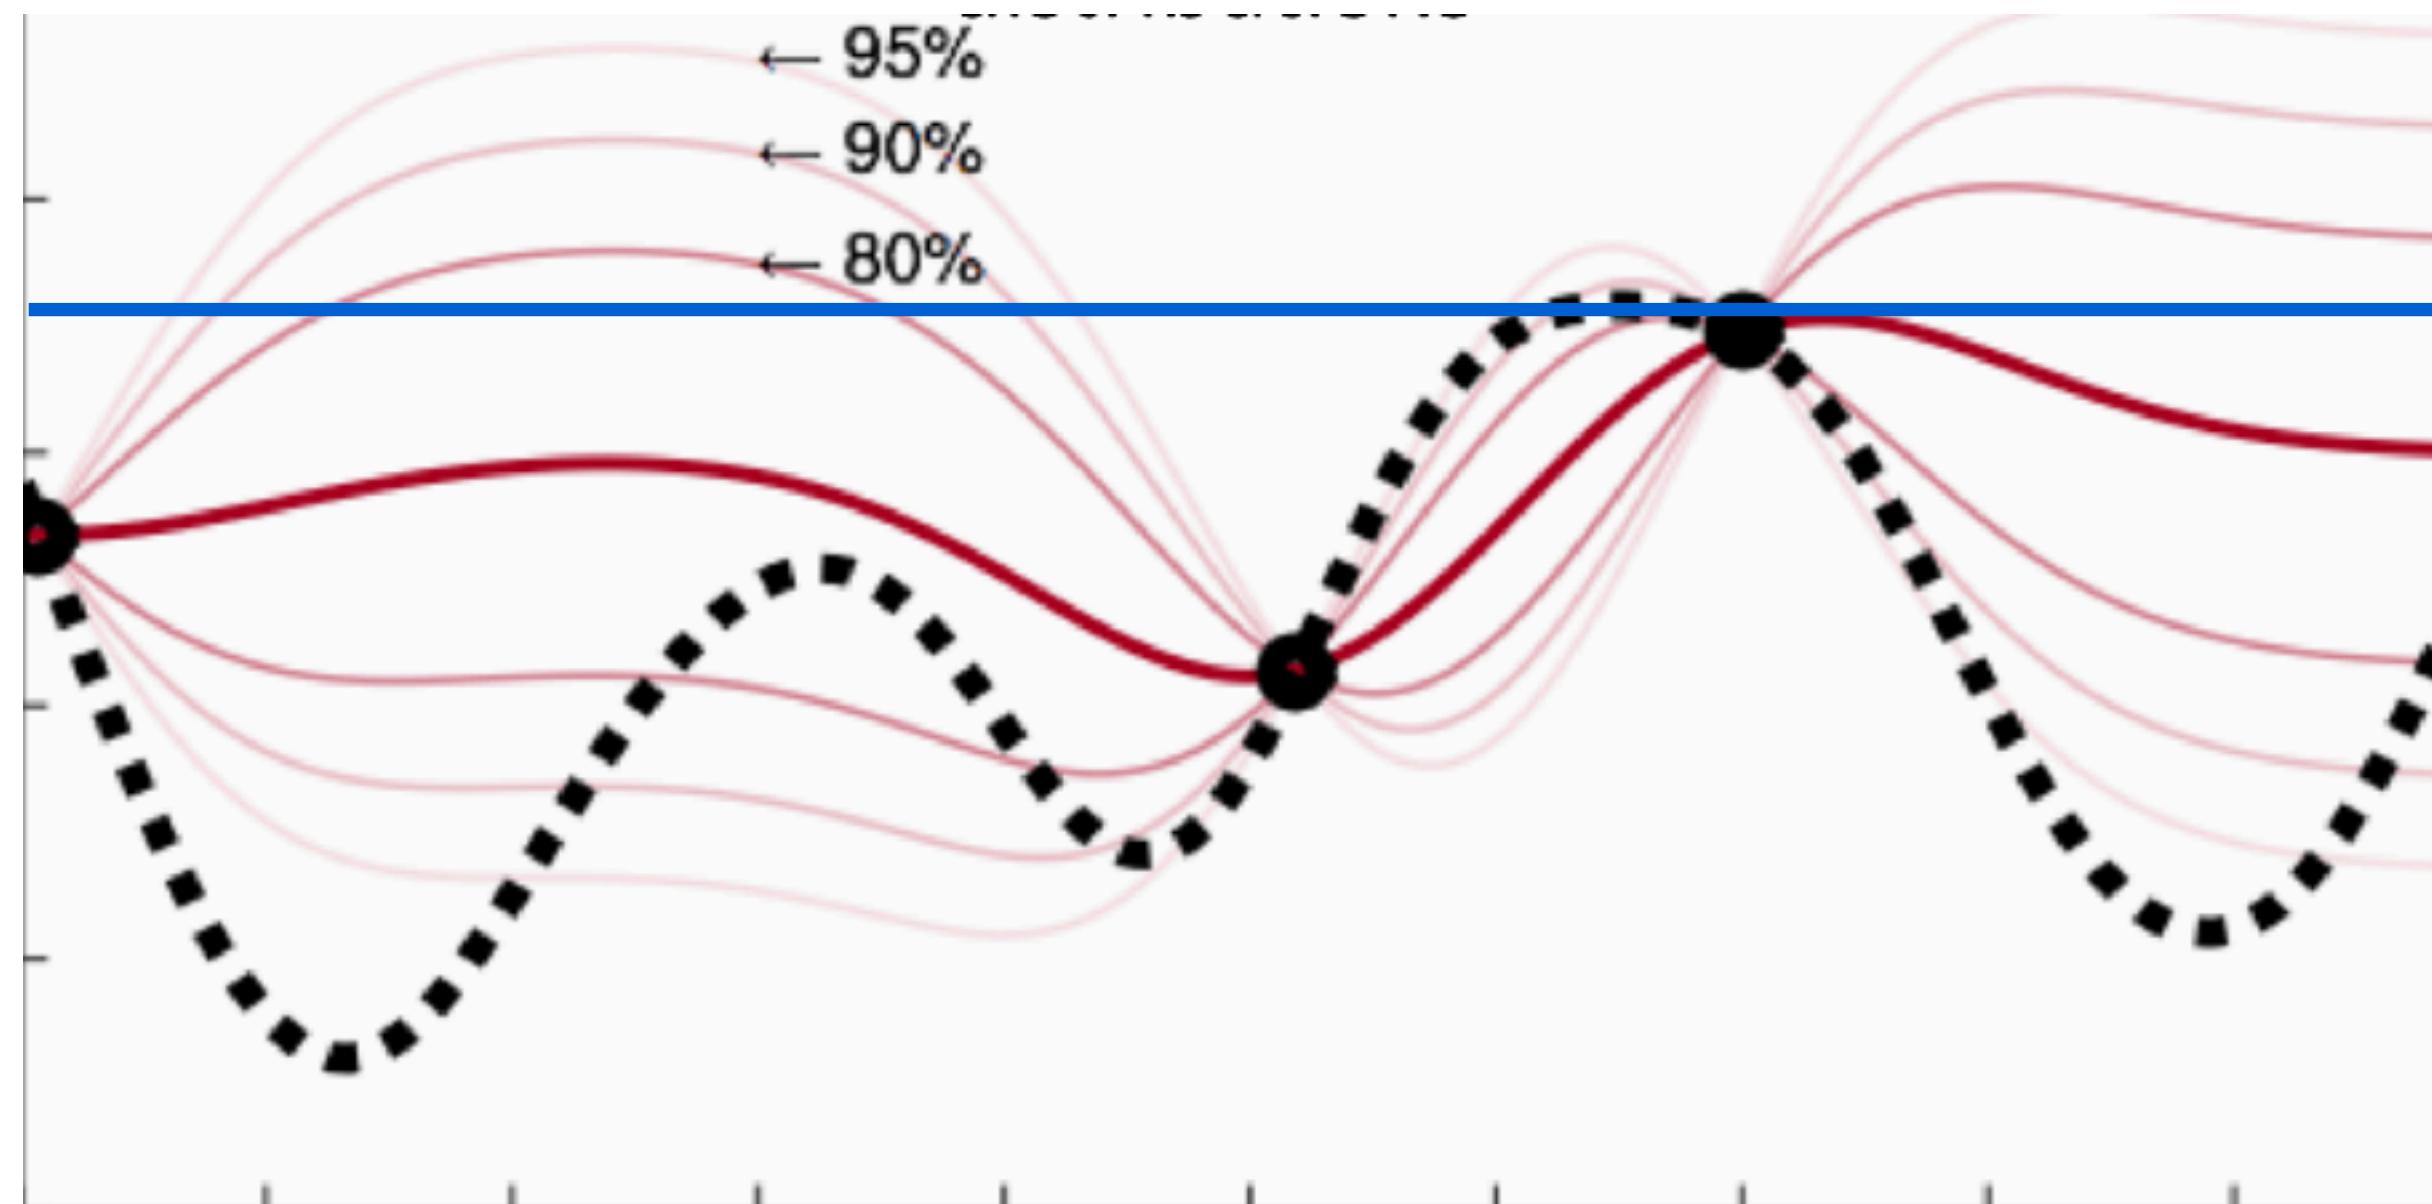
Probability of Improvement



Upper Confidence Bound



Expected Improvement



Acquisition Function

- Thompson Sampling from the posterior GP (TS)
- Probability of Improvement (PI)
- Upper Confidence Bound (UCB)
- **Expected Improvement (EI)**



Maximizing the Acquisition Function

- piecewise infinitely smooth
- gradient-based techniques work
- modified Monte-Carlo techniques are typically used



Optimizing Performance Parameters

```
myExpt = Optimizer.declareDevice(Parm1: {Int, Min1, Max1},  
                                  Parm2: {Real, Min2, Max2},  
                                  Parm3: {Enum, enum1, enum2, enum3}  
                                  ...)  
  
myExpt.setSLA(...)                                // set performance SLA  
myExpt.setTerminationCriteria(...)                // set termination criteria  
  
while (!myExpt.shouldTerminate()) {  
    parmSuggestion = myExpt.suggest()              // get another test suggestion  
    newRun = myDevice.test(parmSuggestion)          // test device at given setting  
    if (myExpt.isValid(newRun)) {  
        myExpt.update(parmSuggestion, newRun)        // is SLA met?  
    }  
}  
  
return myExpt.bestConfig()
```



Bayesian Optimization

```
myExpt = Optimizer.declareDevice(Parm1: {Int, Min1, Max1},  
                                 Parm2: {Real, Min2, Max2},  
                                 Parm3: {Enum, enum1, enum2, enum3}  
                                 ...)  
  
myExpt.setSLA(...)                                // set performance SLA  
myExpt.setTerminationCriteria(...)                // set termination criteria  
  
while (!myExpt.shouldTerminate()) {  
    parmSuggestion = myExpt.suggest()              // get another test suggestion  
    newRun = myDevice.test(parmSuggestion)          // test device at given setting  
    if (myExpt.isValid(newRun)) {  
        myExpt.update(parmSuggestion, newRun)        // is SLA met?  
    }  
}  
  
return myExpt.bestConfig()
```

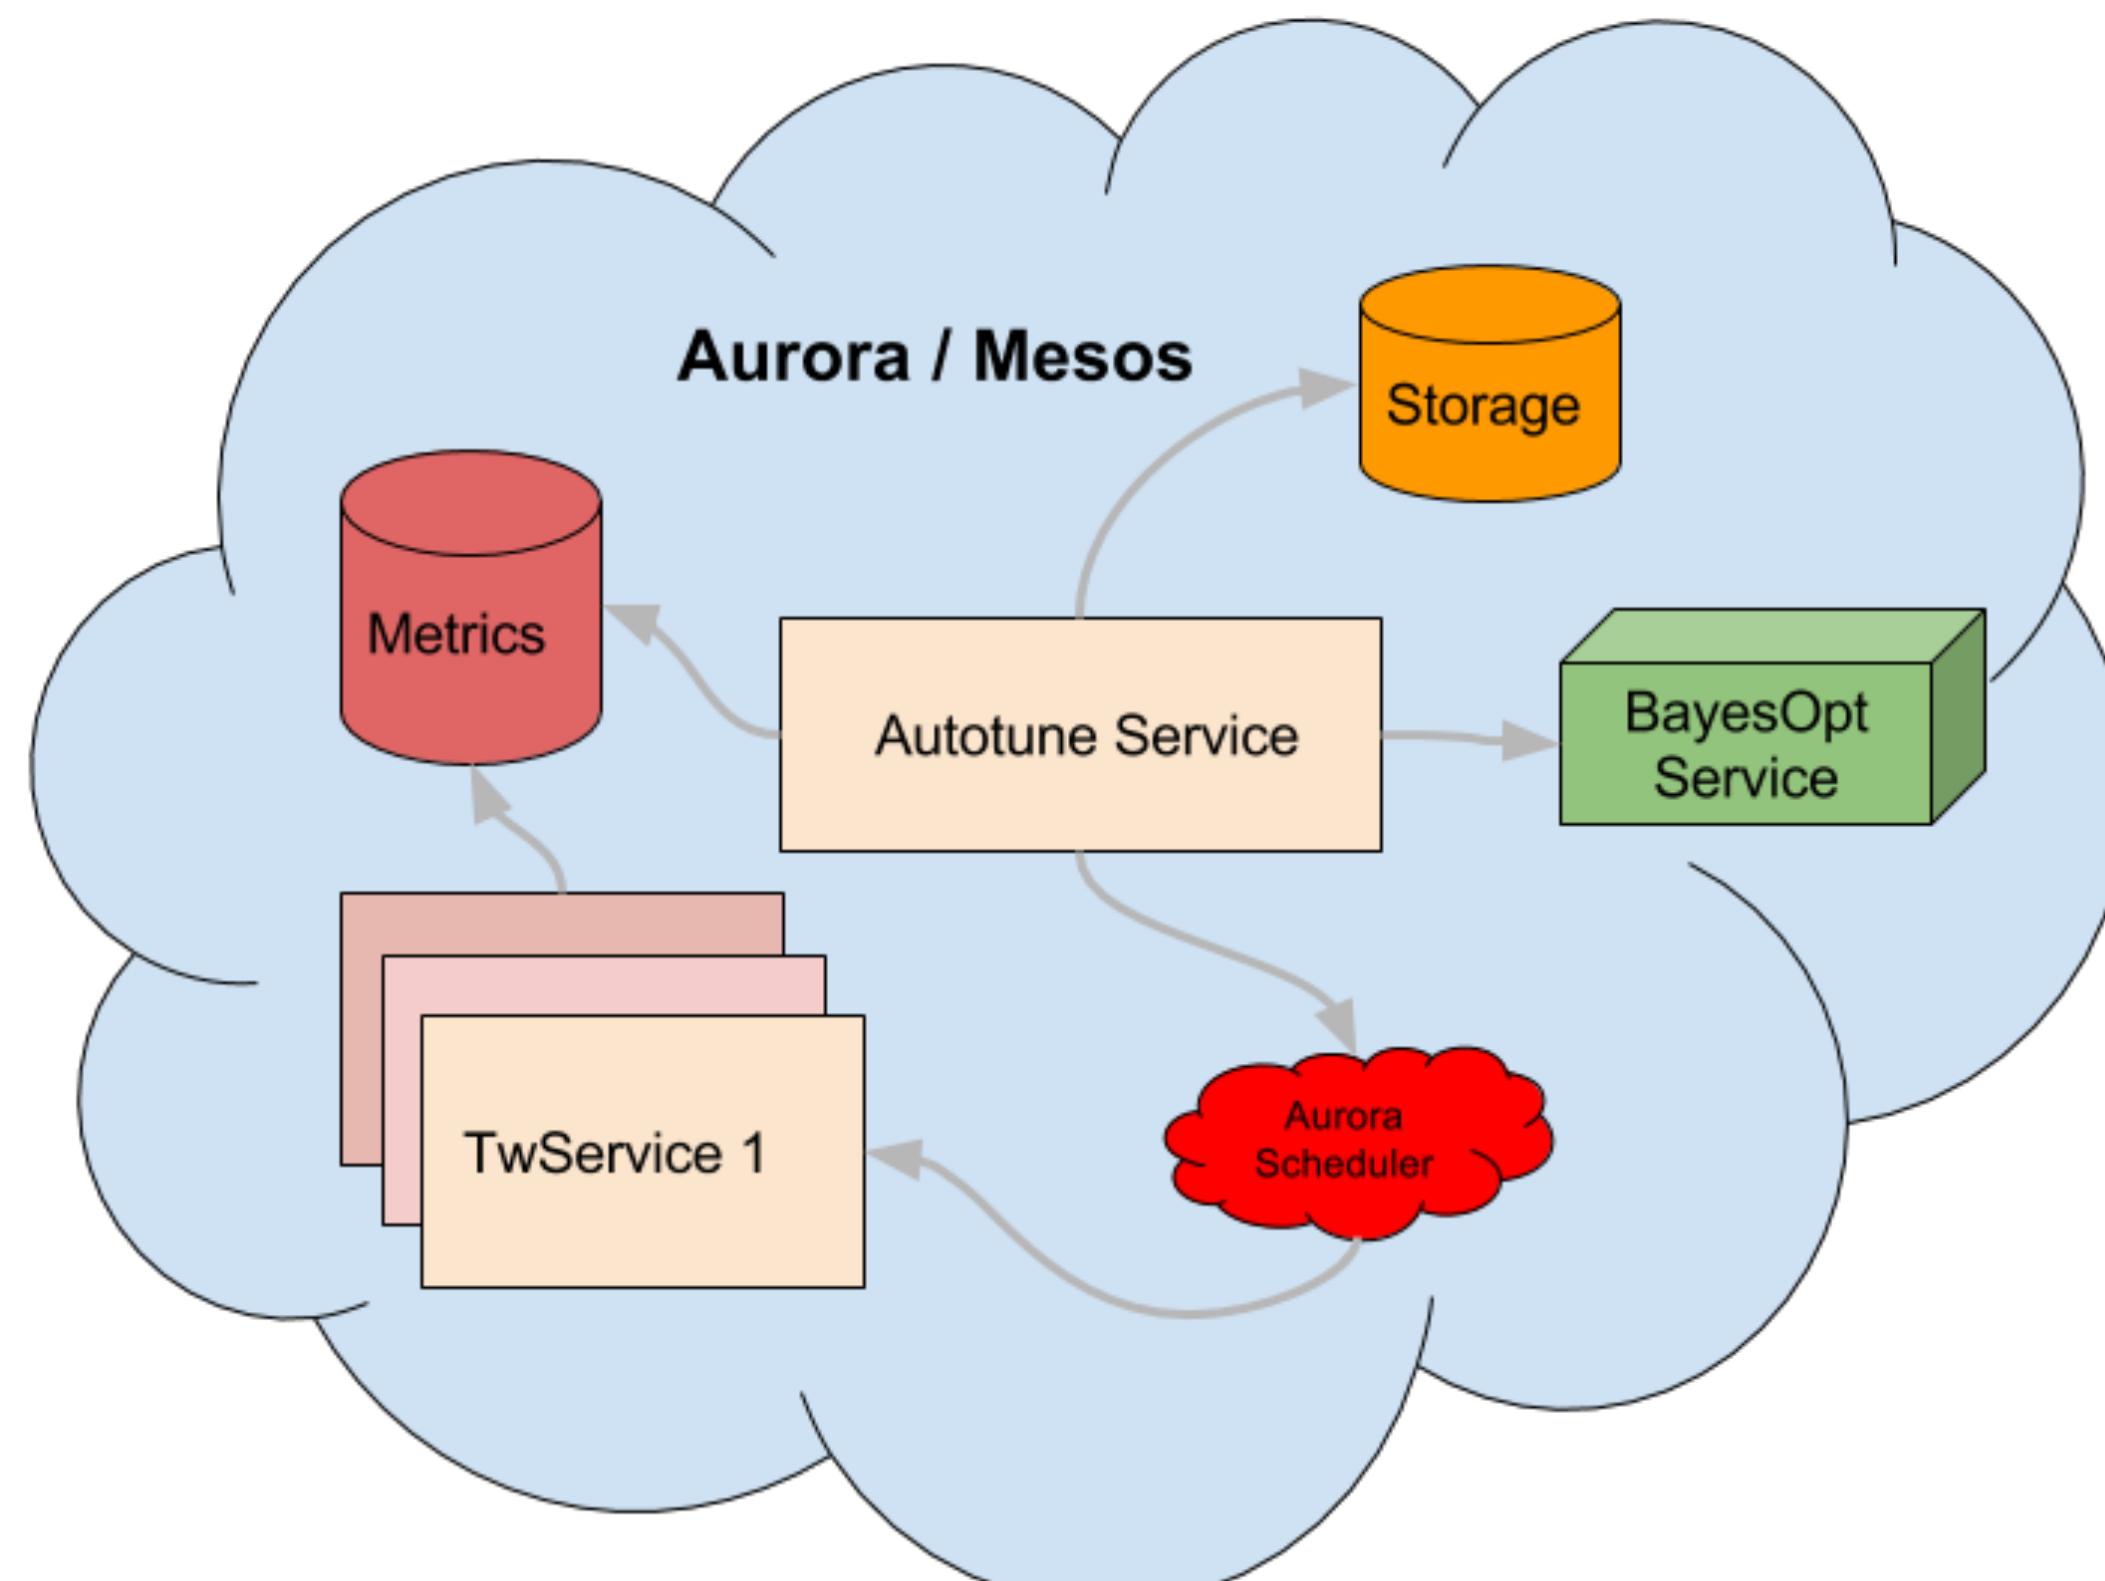


Constraints

```
myExpt = Optimizer.declareDevice(Parm1: {Int, Min1, Max1},  
                                  Parm2: {Real, Min2, Max2},  
                                  Parm3: {Enum, enum1, enum2, enum3}  
                                  ...)  
  
myExpt.setSLA(...)                                // set performance SLA  
myExpt.setTerminationCriteria(...)                // set termination criteria  
  
while (!myExpt.shouldTerminate()) {  
    parmSuggestion = myExpt.suggest()              // get another test suggestion  
    newRun = myDevice.test(parmSuggestion)          // test device at given setting  
    if (myExpt.isValid(newRun)) {  
        myExpt.update(parmSuggestion, newRun)       // is SLA met?  
    }  
}  
  
return myExpt.bestConfig()
```

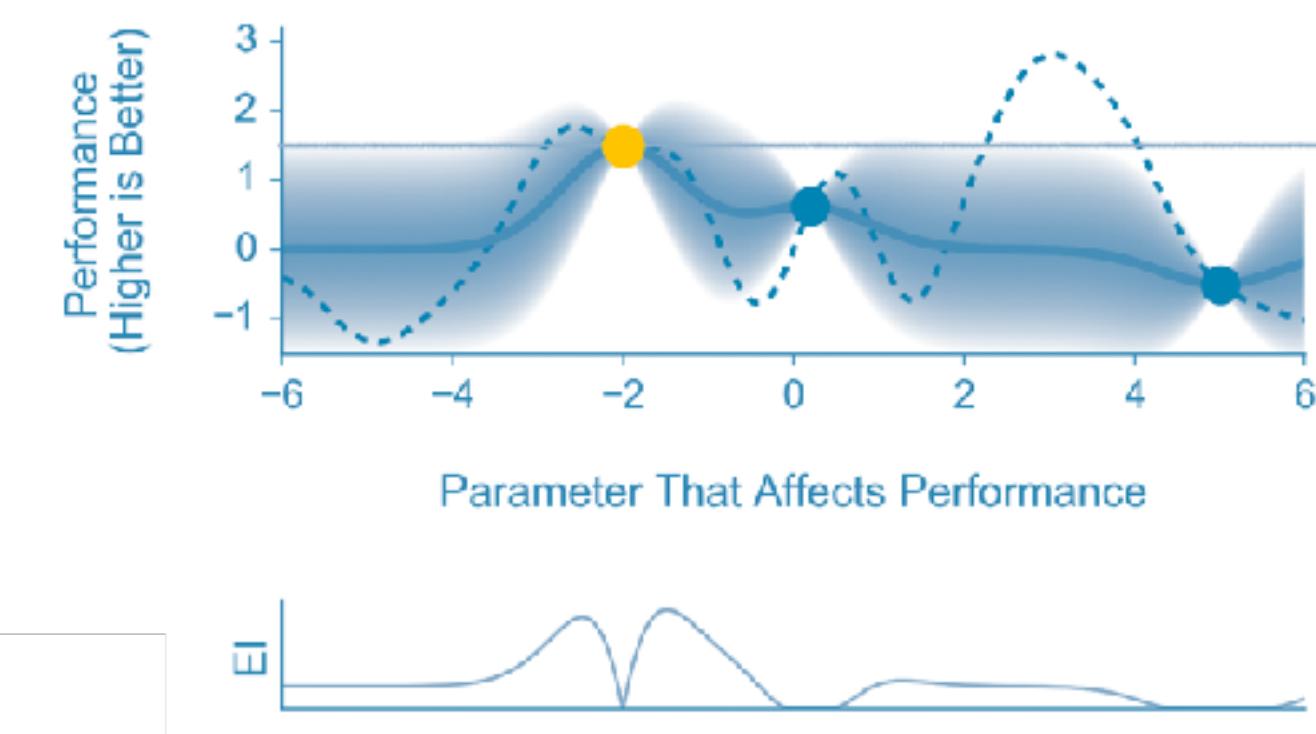


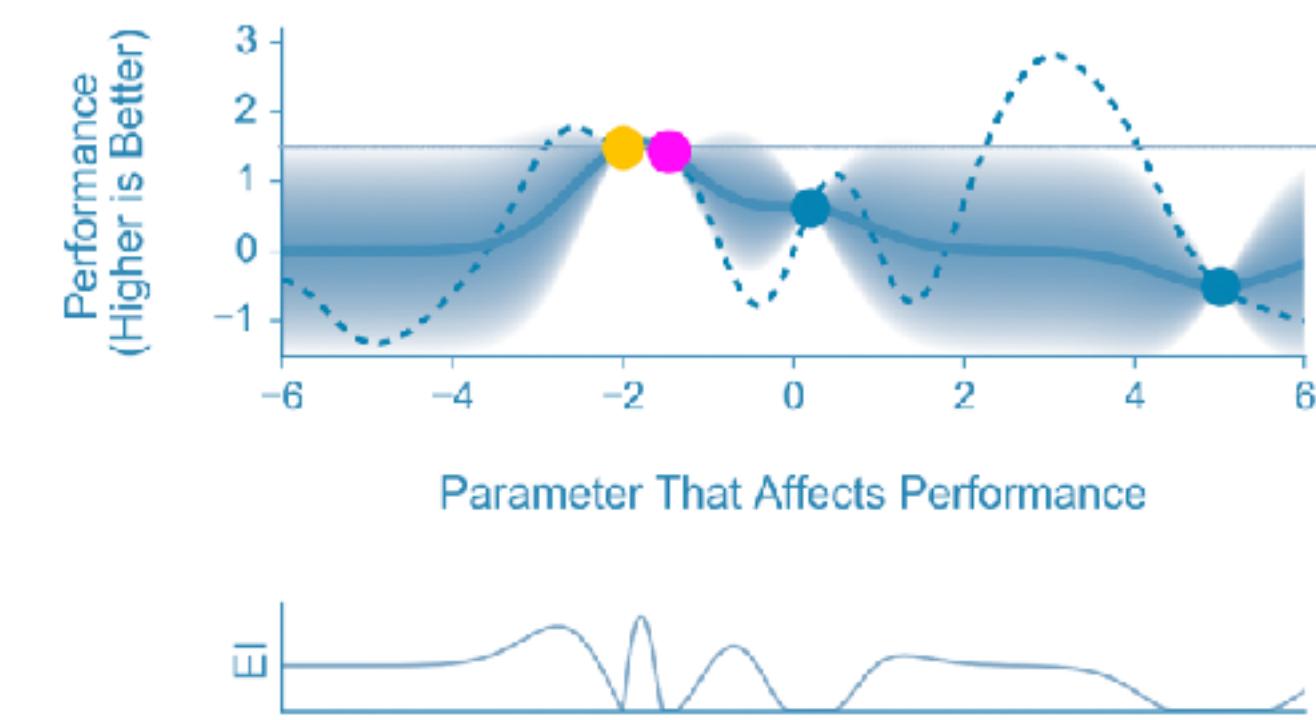
AUTOTUNE AS A SERVICE

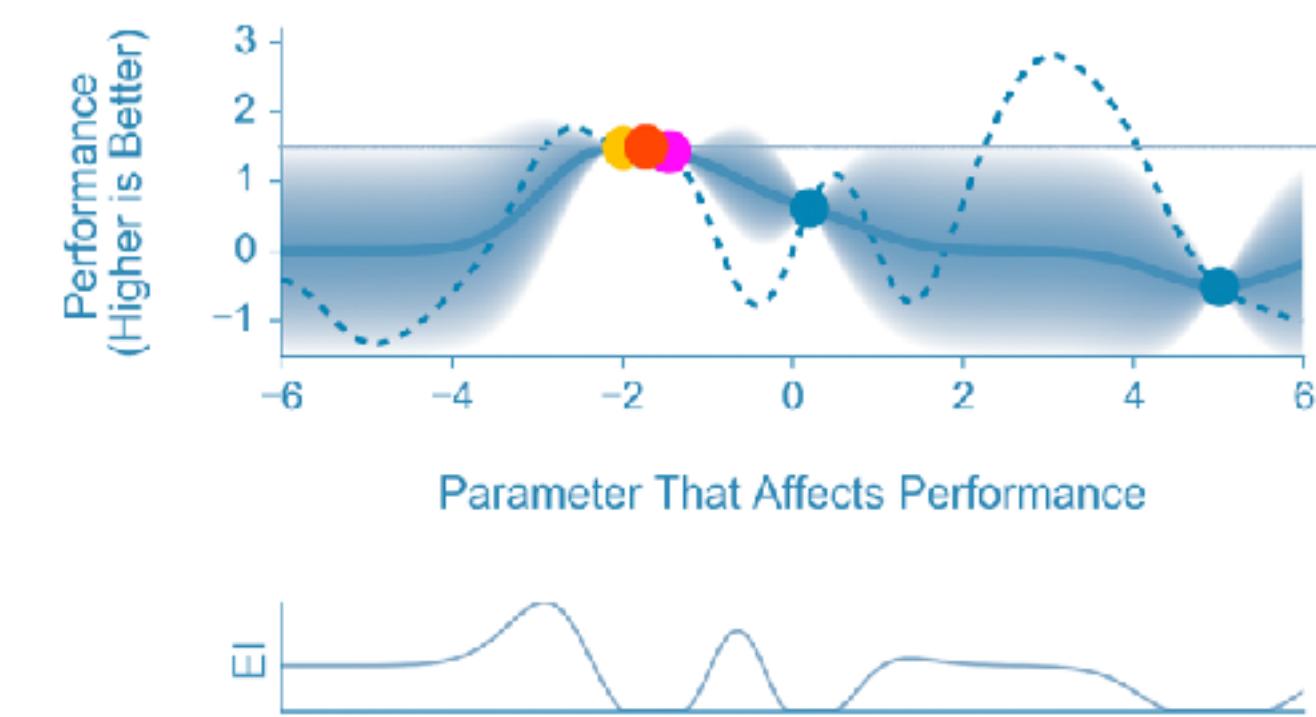


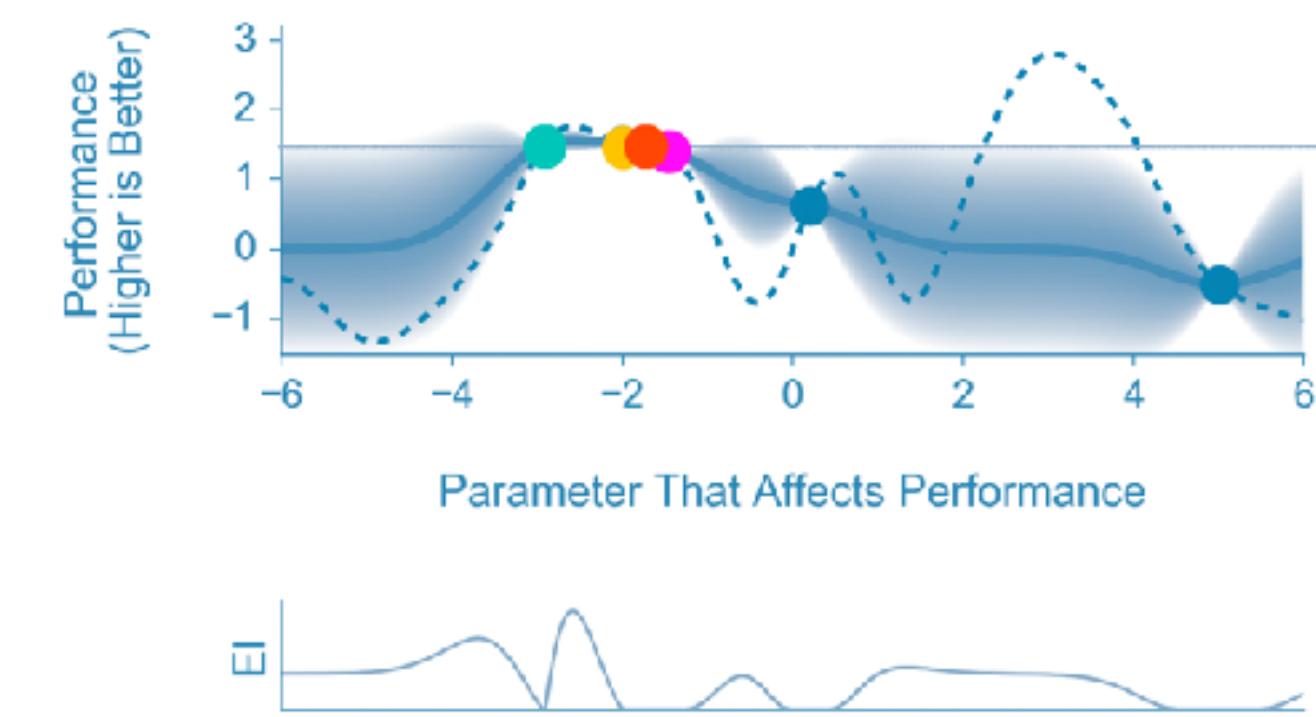


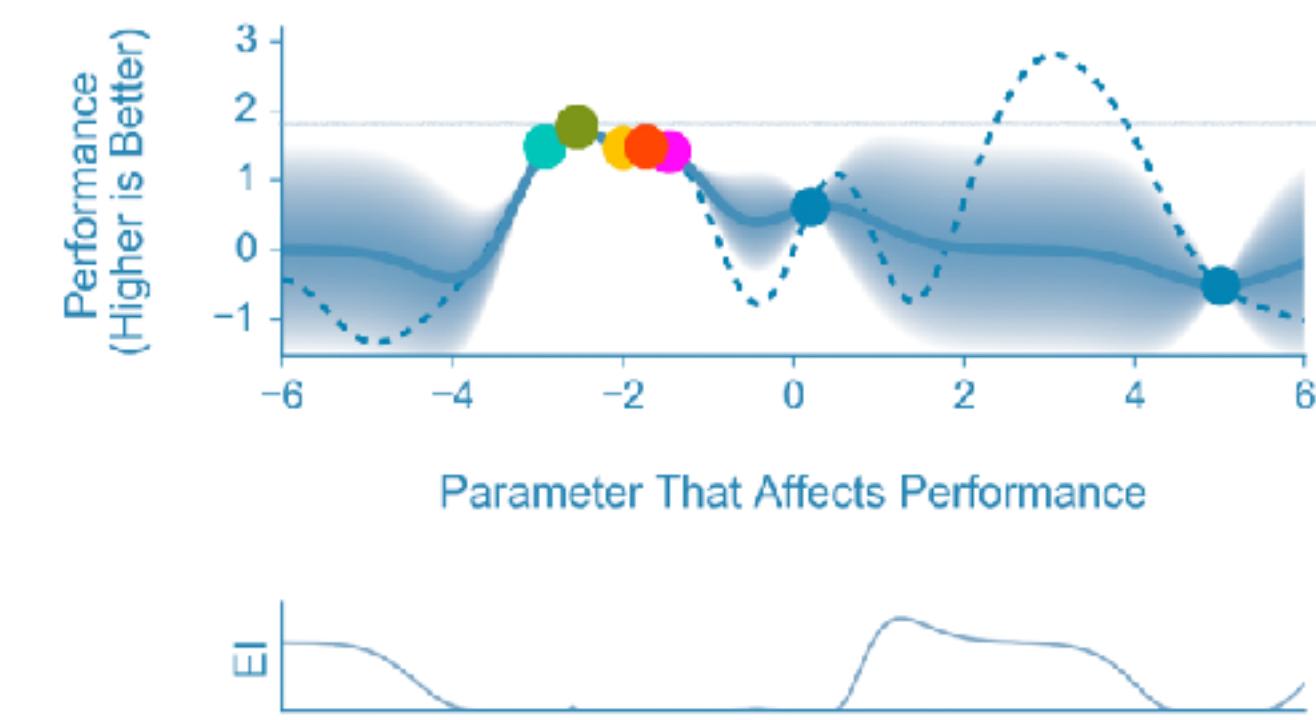


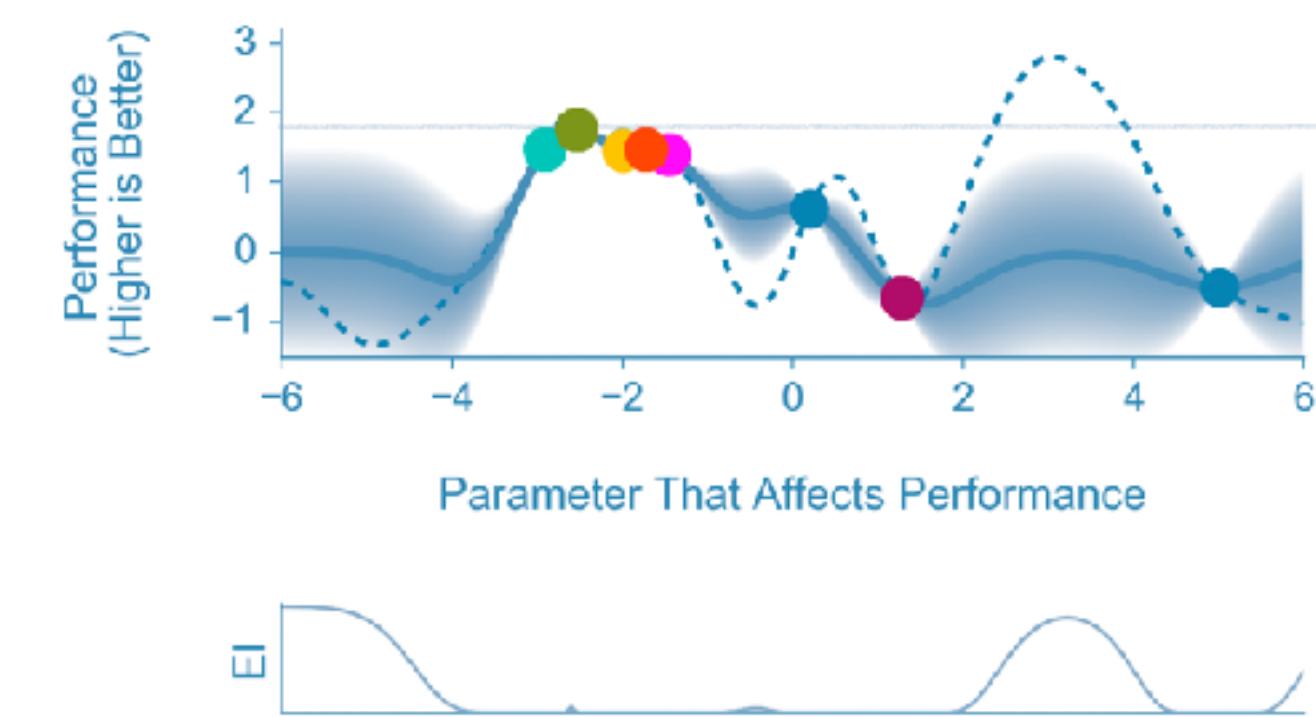


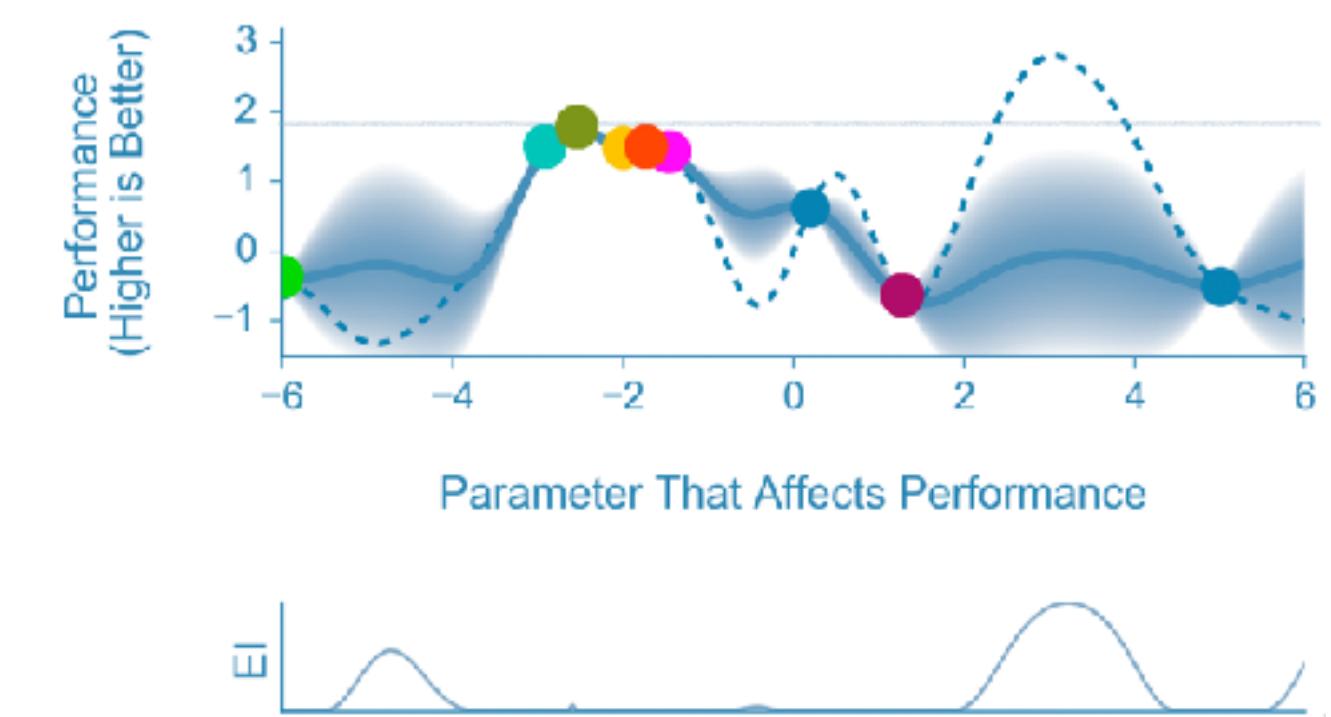


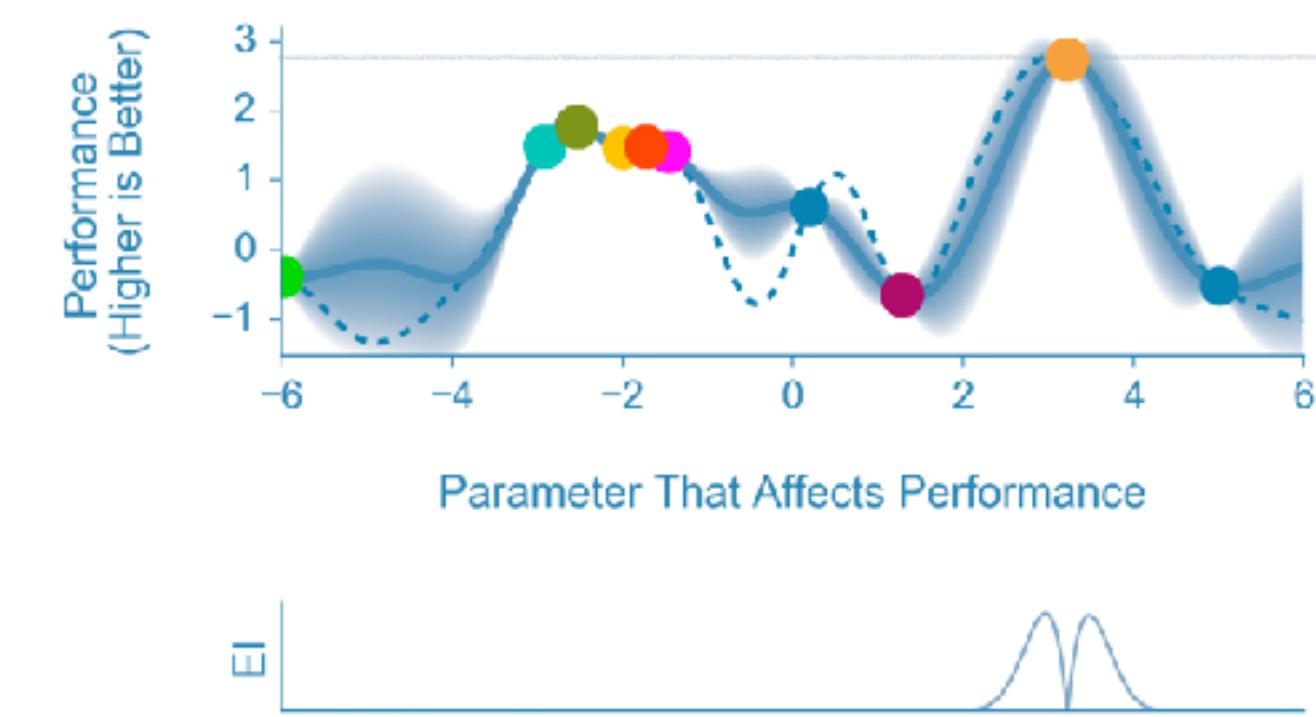


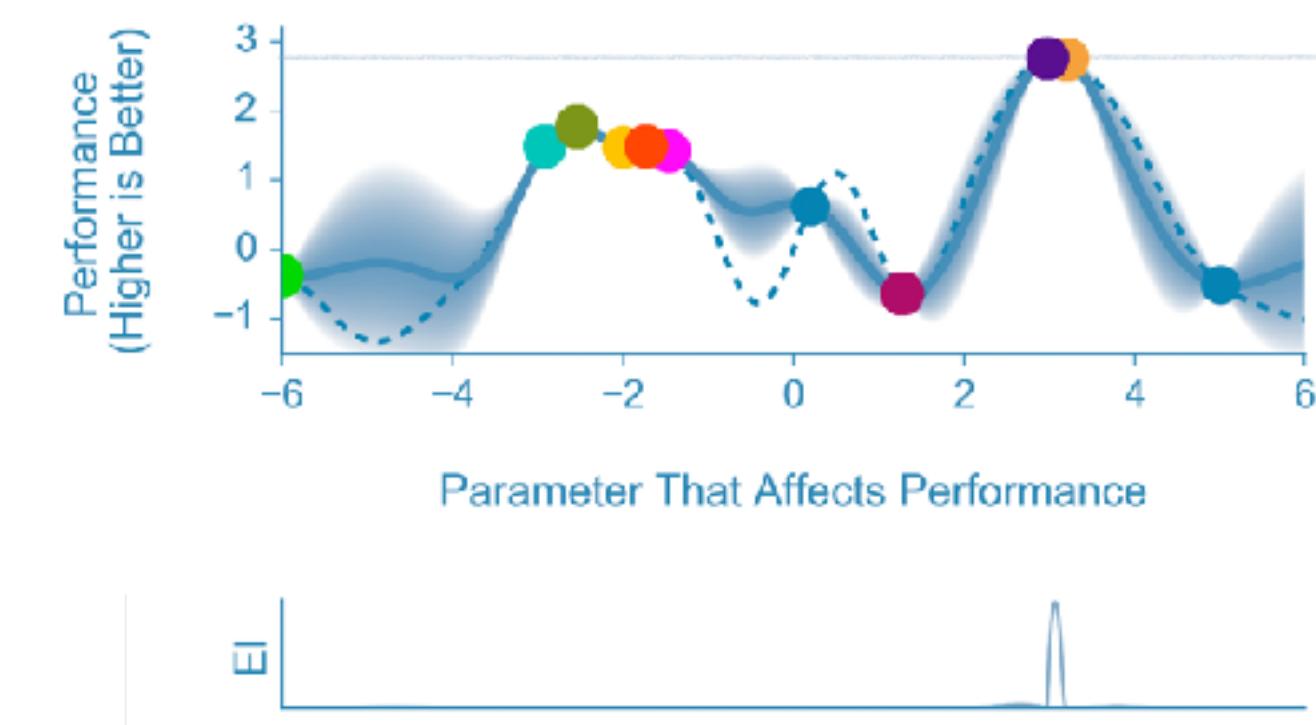


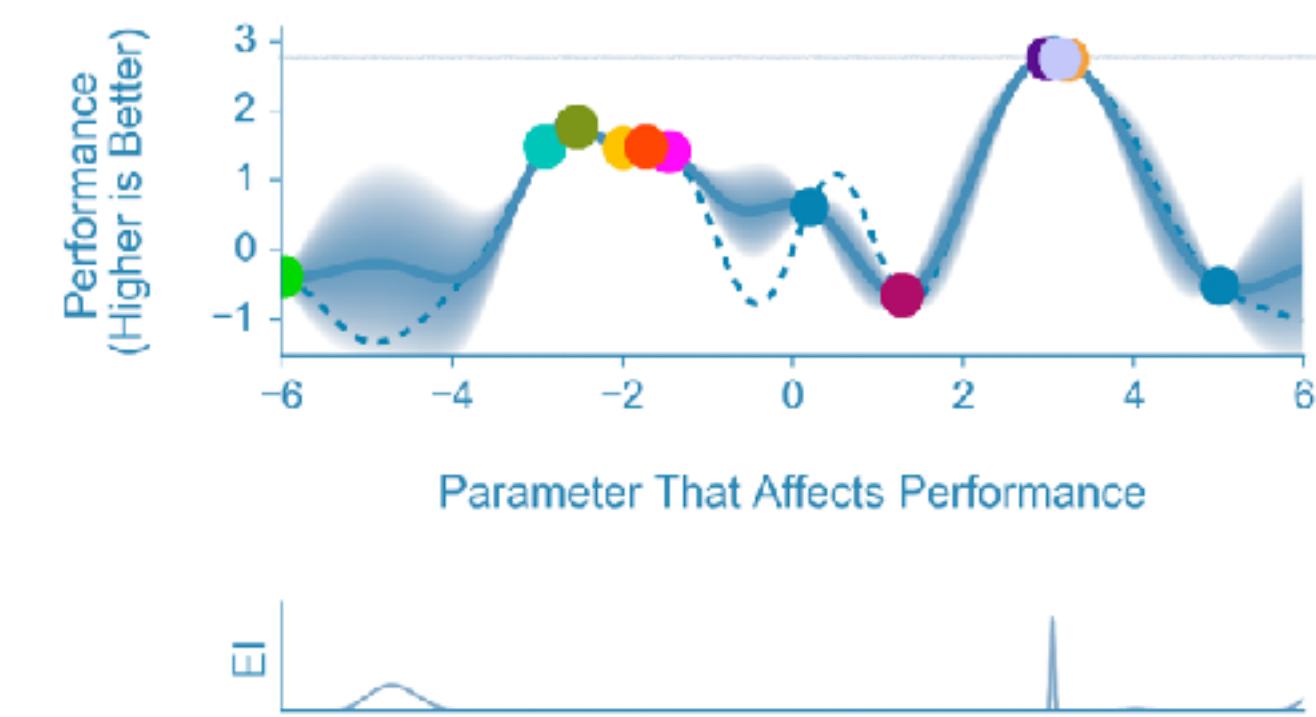


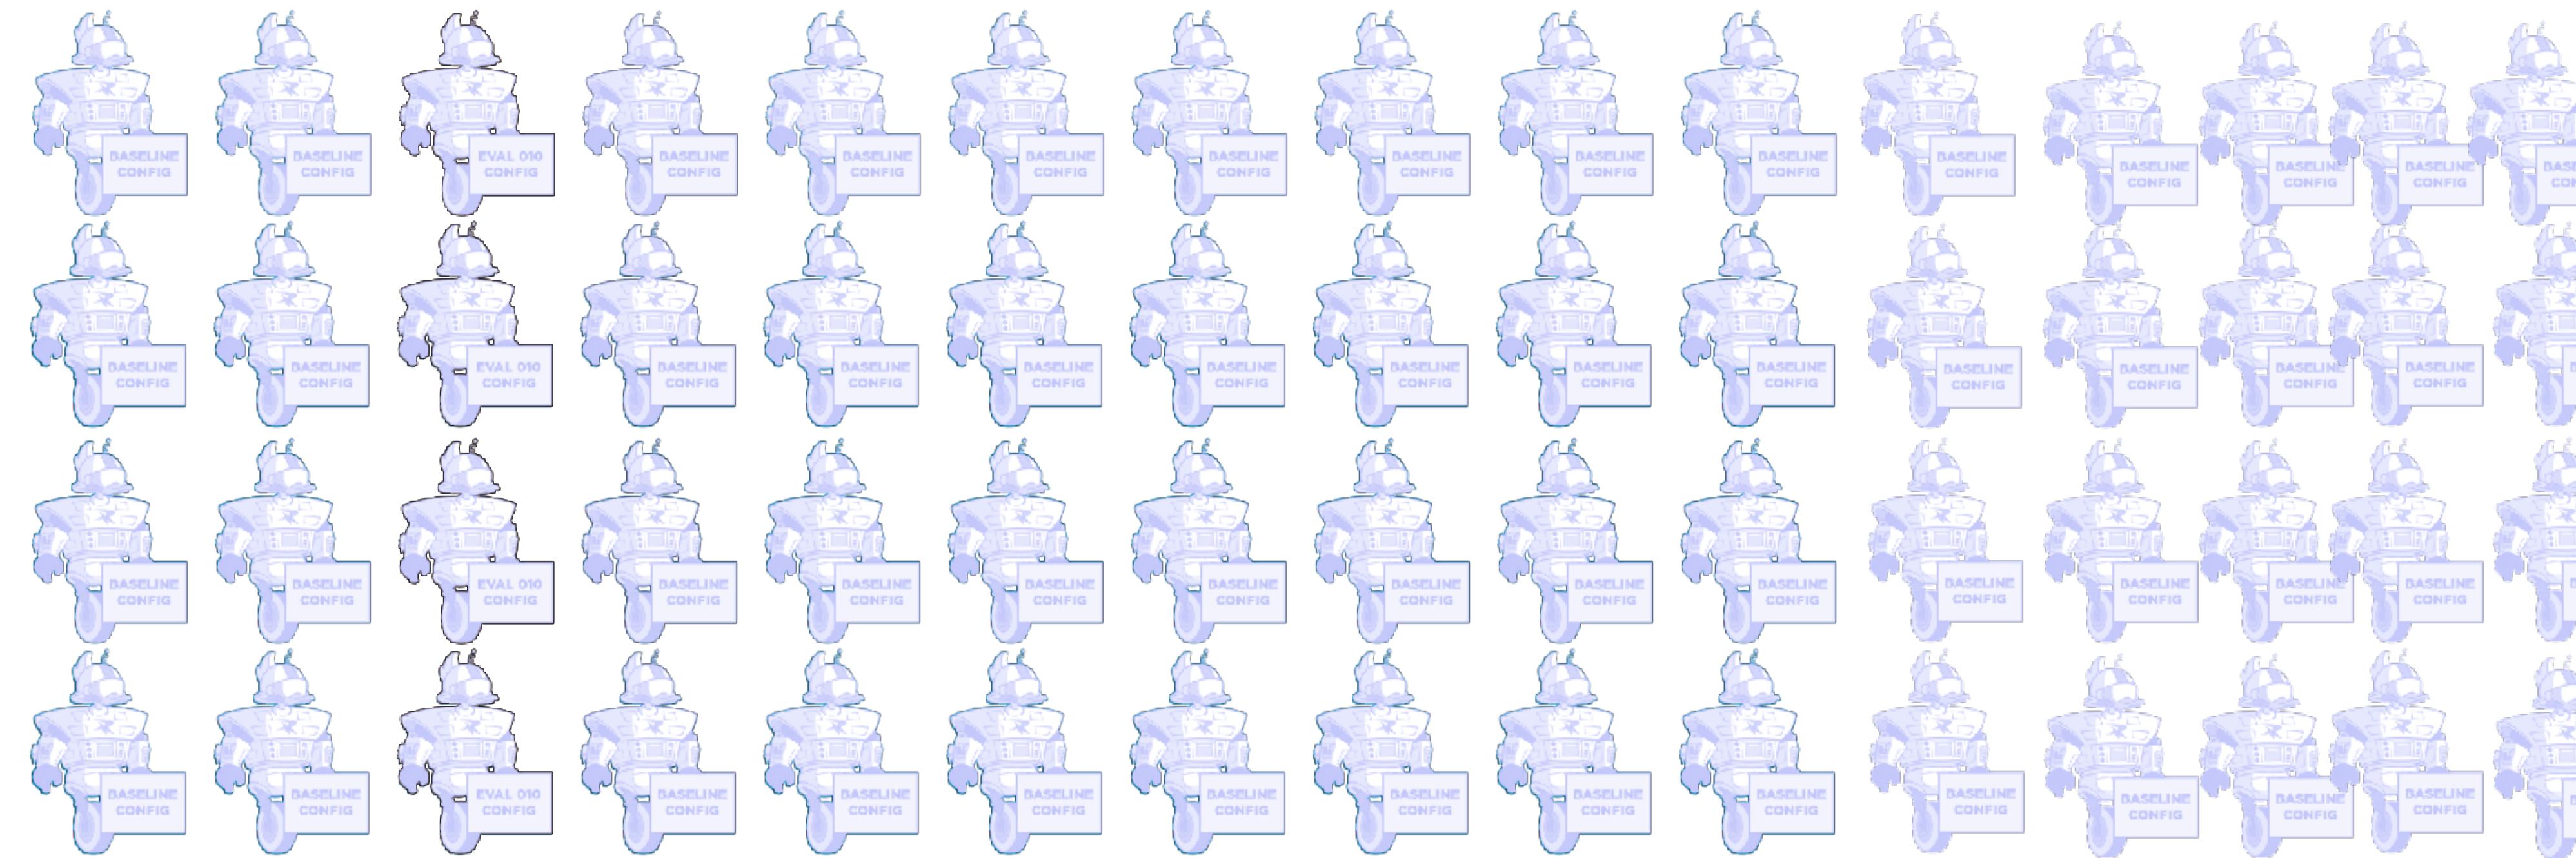
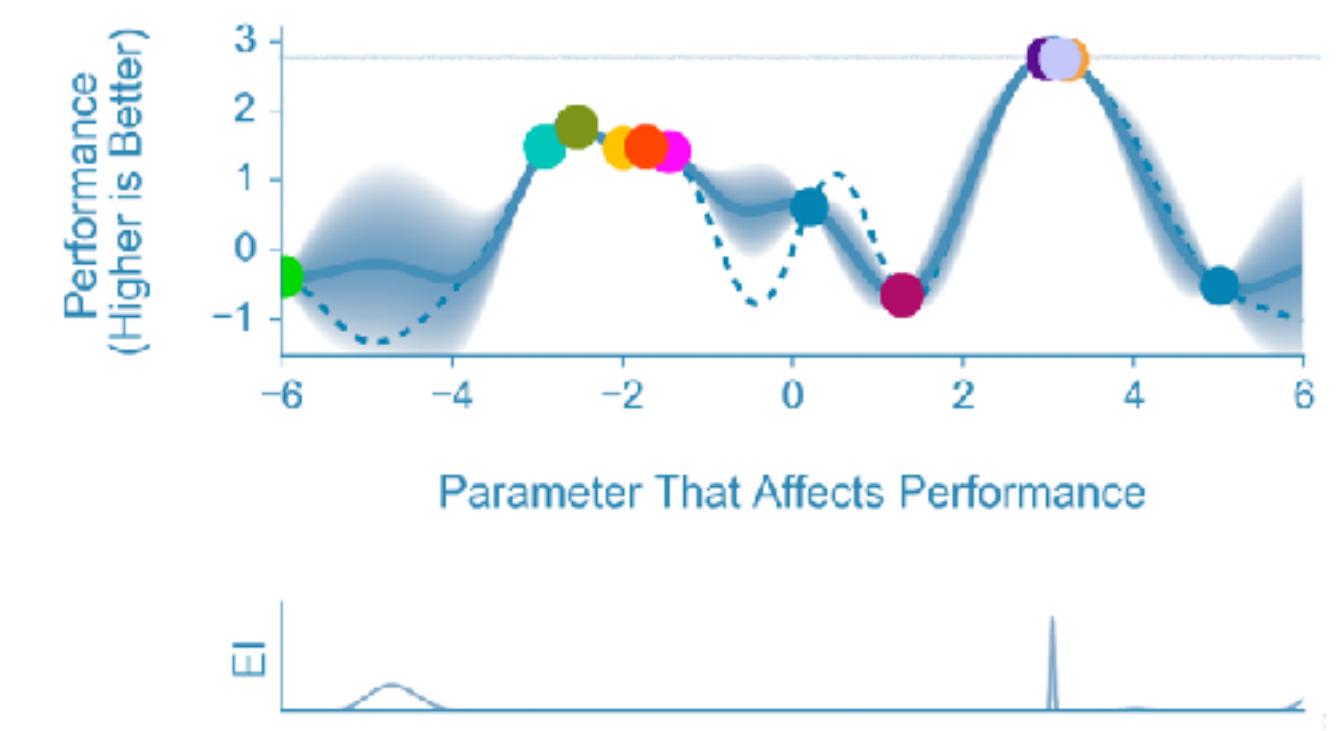




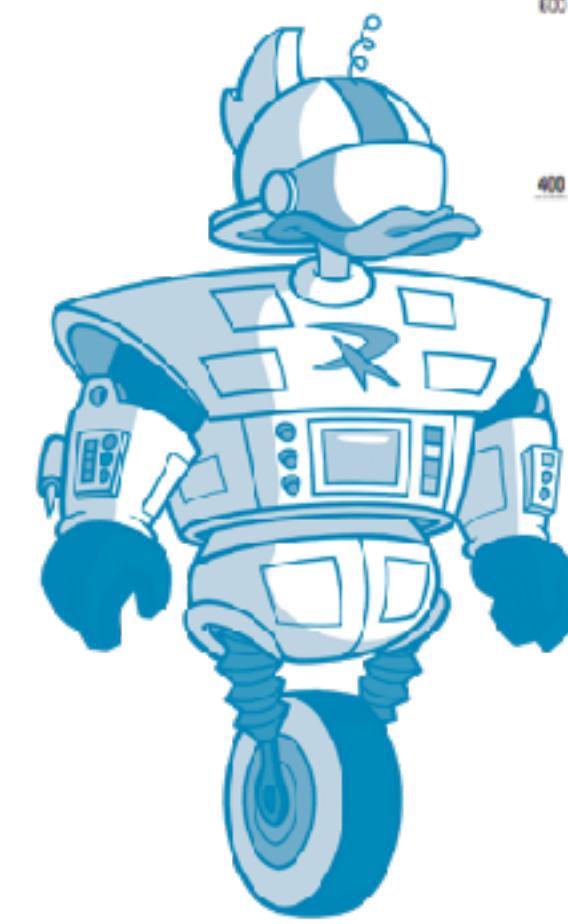
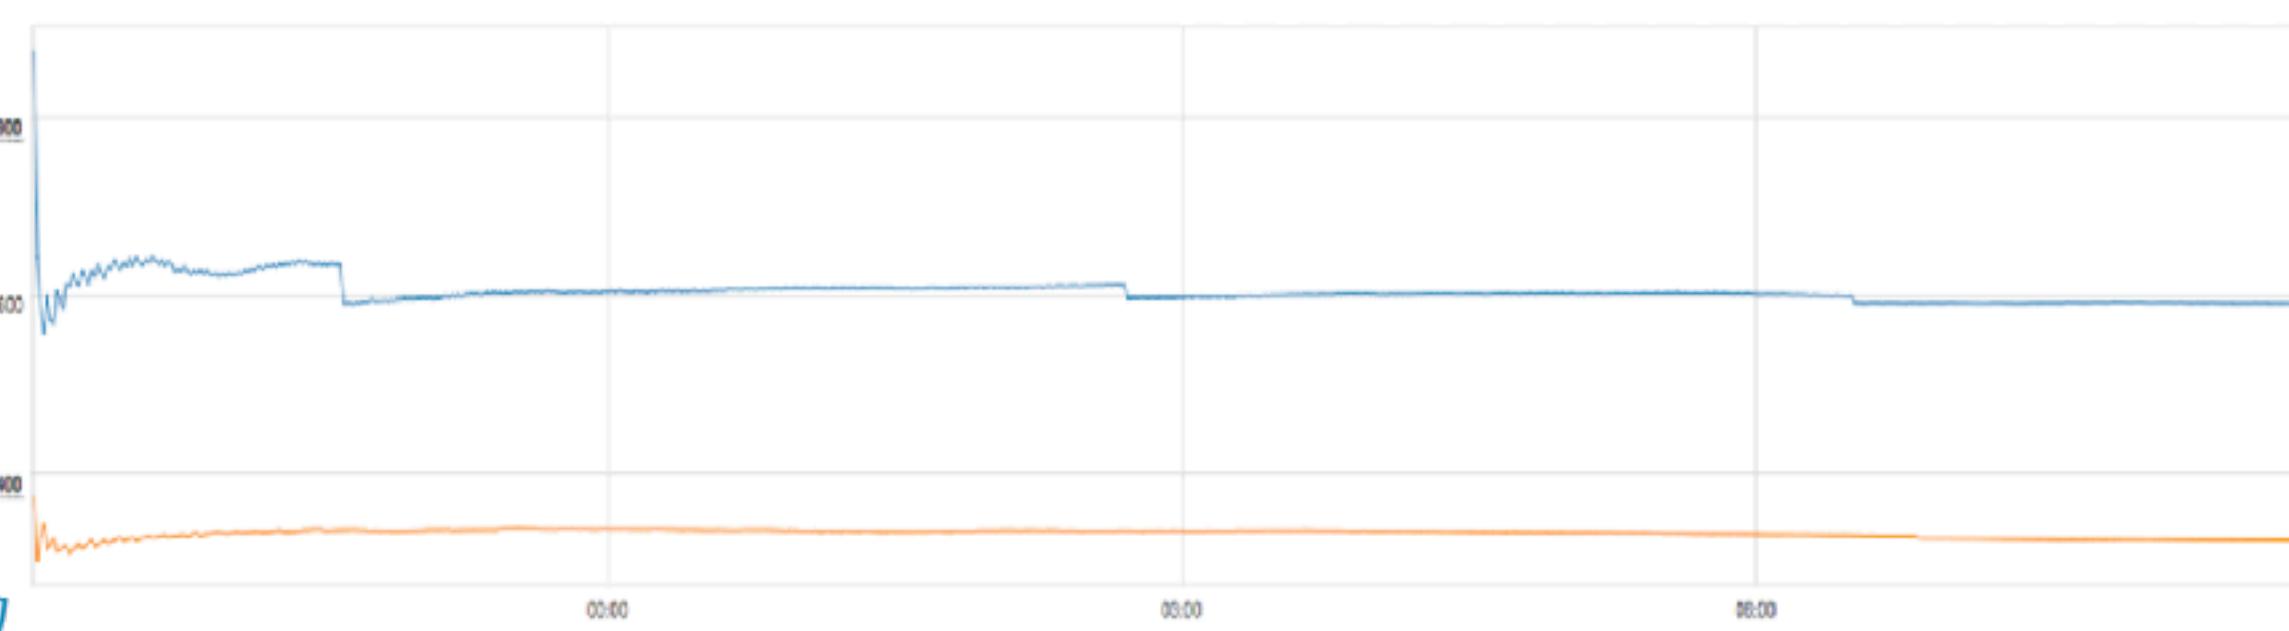
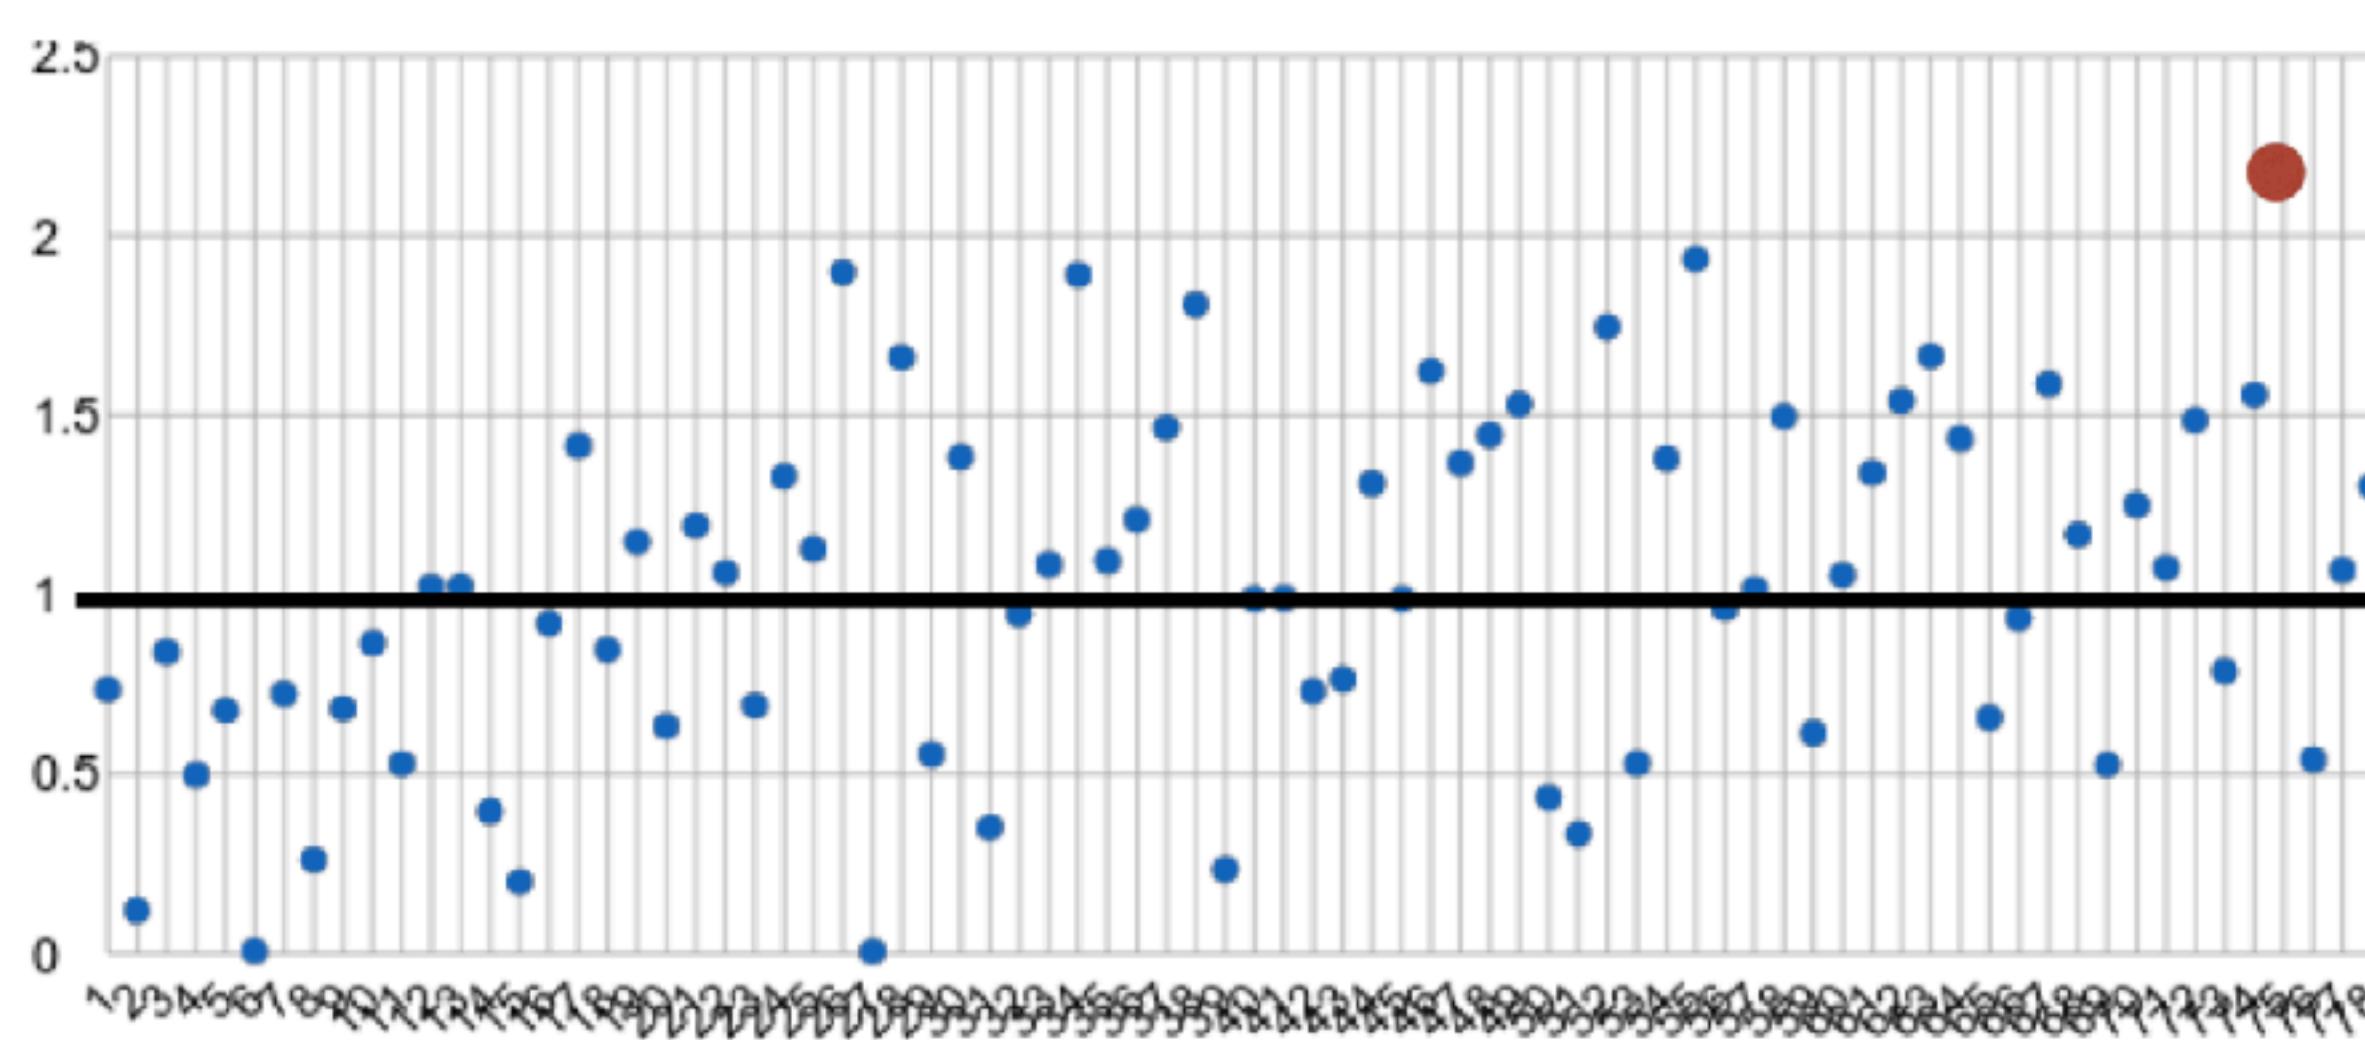








GizmoDuck & Garbage Collection Overhead via Tuning JVM Parameters



TweetyPie & CPU Utilization via Tuning Graal JIT Parameters

